


Linux Power Management – Challenges in Modern Servers and how SmartNICs can help

Marco Liess, Andreas Herkersdorf
Technical University of Munich
Chair of Integrated Systems
Berlin, November 13th 2025



About Me



Marco Liess (✉ marco.liess@tum.de)

PhD Candidate @ Chair of Integrated Systems
with Prof. Andreas Herkersdorf



M. Sc. in Electrical and Computer Engineering @ TUM




Research on energy-efficient and adaptive SmartNICs (“intelligent network interfaces”) for 6G RAN/Edge Nodes



Teaching assistant for Chip Multicore Processors (EI7271)


Architectures

- MPSOC platforms
- Hardware accelerators
- 2D/3D NoC interconnect
- Memory hierarchies



Methods

- Cross-layer resilience
- Bio-inspired reinforcement ML
- Power management
- Diagnosis on Chip



Application Areas



- Networking
- Automotive
- Visual Computing
- Near Memory Acceleration

Prototyping Lab



- FPGA-based emulation
- Network testing
- Spectrum Analyzer
- x86-based Manycore

Importance of Energy-efficiency

Energy Production and Compute Consumption





Figure: Worldwide Energy Production and Compute Consumption ¹

¹ Semiconductor Research Corporation. *The Decadal Plan for Semiconductors 2021*. 2021

Importance of Energy-efficiency

Global Data Center Electricity Consumption ¹




¹ IEA (2025), *Energy and AI*, IEA, Paris <https://www.iea.org/reports/energy-and-ai>, Licence: CC BY 4.0

Importance of Energy-efficiency

How can it be improved (from a Hardware perspective)?

Importance of Energy-efficiency


How can it be improved (from a Hardware perspective)?



Specialized Architectures
(e.g. Tensor Processing Unit) → *High Ops/Watt*

Importance of Energy-efficiency

How can it be improved (from a Hardware perspective)?



Specialized Architectures
(e.g. Tensor Processing Unit)


→ *High Ops/Watt*



Heterogeneous Architectures
("Choosing the right tool for the right job")

Importance of Energy-efficiency

How can it be improved (from a Hardware perspective)?



Specialized Architectures
(e.g. Tensor Processing Unit) → *High Ops/Watt*



Heterogeneous Architectures
„*Choosing the right tool for the right job*“)



Intelligent Resource Management
(Power management, scheduling, task distribution, etc...)

Overview


1. CPU Power Management – Linux CPUFreq
2. Analytic Considerations of Asynchronous Traffic-driven DVFS
3. EcoNIC: SmartNIC-assisted Linux Power Management
4. Conclusion & Outlook

CPU Runtime Power Management

Between Running, Walking, Dozing and Sleeping

Runtime Configuration Options:

- Frequency Adjustments
 - Main driver is dynamic power consumption on clock ticks
 - Less clock ticks → less power
 - Further enables voltage reduction → nonlinear savings
- „Sleeping“
 - „Clock gating“ → disable parts of the CPU
 - “Turn off“ supply voltage



CPU Runtime Power Management

Sleep states – C-States

→ CPU Architecture supports certain C-States

→ C-State activation / selection managed by OS

e.g. 350 µs on
Ryzen 5 5600G




https://www.thomas-krenn.com/en/wiki/Processor_P-states_and_C-states

CPU Runtime Power Management

Dynamic Voltage Frequency Scaling (DVFS) – P-States

- P-State availability and characteristics (transition latency, frequencies, etc.) depend on CPU architecture
- P-State management can be performed by OS or CPU itself (e.g. Intel SpeedShift, AMD CPPC)



<https://www.kernel.org/doc/html/v6.0/admin-guide/pm/amd-pstate.html>

DVFS – P-State Management

Power consumption of digital circuits

- 1) Dynamic power consumption:

$$P_{dyn} = \alpha C \cdot f_{clk} \cdot V_{dd}^2$$

α : activity, C: capacity,
 f_{clk} : clock frequency, V_{dd} : supply voltage




Figure: Power consumption of one core of a Ryzen 5 5600G 6-core CPU

DVFS – P-State Management

Power consumption of digital circuits

- 1) Dynamic power consumption:

$$P_{dyn} = \alpha C \cdot f_{clk} \cdot V_{dd}^2$$

α : activity, C: capacity,

f_{clk} : clock frequency, V_{dd} : supply voltage

- 2) Available processing capacity:

$$MIPS_{prov} = N \cdot f_{clk} \cdot IPC$$

N : # of cores, f_{clk} : clock frequency, IPC : instructions per cycle




Figure: Power consumption of one core of a Ryzen 5 5600G 6-core CPU

DVFS – P-State Management

Power consumption of digital circuits

- 1) Dynamic power consumption:

$$P_{dyn} = \alpha C \cdot f_{clk} \cdot V_{dd}^2$$

α : activity, C: capacity,
 f_{clk} : clock frequency, V_{dd} : supply voltage

- 2) Available processing capacity:

$$MIPS_{prov} = N \cdot f_{clk} \cdot IPC$$

N : # of cores, f_{clk} : clock frequency, IPC : instructions per cycle

- 3) Required processing capacity

$$MIPS_{req}$$




Figure: Power consumption of one core of a Ryzen 5 5600G 6-core CPU

DVFS – P-State Management

Power consumption of digital circuits

- 1) Dynamic power consumption:

$$P_{dyn} = \alpha C \cdot f_{clk} \cdot V_{dd}^2$$

α : activity, C: capacity,
 f_{clk} : clock frequency, V_{dd} : supply voltage

- 2) Available processing capacity:

$$MIPS_{prov} = N \cdot f_{clk} \cdot IPC$$

N: # of cores, f_{clk} : clock frequency, IPC: instructions per cycle

- 3) Required processing capacity

$$MIPS_{req}$$




Figure: Power consumption of one core of a Ryzen 5 5600G 6-core CPU

→ Use f_{clk} to match supply to demand!

DVFS – P-State Management

Power consumption of digital circuits

- 1) Dynamic power consumption:

$$P_{dyn} = \alpha C \cdot f_{clk} \cdot V_{dd}^2$$

α : activity, C: capacity,
 f_{clk} : clock frequency, V_{dd} : supply voltage

- 2) Available processing capacity:

$$MIPS_{prov} = N \cdot f_{clk} \cdot IPC$$

N: # of cores, f_{clk} : clock frequency, IPC: instructions per cycle

- 3) Required processing capacity

$$MIPS_{req} = ?$$

Depends on application and can be hard to predict!




Figure: Power consumption of one core of a Ryzen 5 5600G 6-core CPU


→ Use f_{clk} to match supply to demand!

DVFS – P-State Management

Linux CPUFreq Subsystem

Basics:

- CPUFreq Core is abstraction layer
- Driver implements CPU Hardware features
- Governor makes decisions



DVFS – P-State Management


Linux CPUFreq Subsystem

Basics:

- CPUFreq Core is abstraction layer
- Driver implements CPU Hardware features
- Governor makes decisions

Popular Linux Governors:

- performance always highest frequency
- powersave always lowest frequency
- ondemand based on current system load
- schedutil based on scheduler's PELT



DVFS – P-State Management

Linux CPUFreq Subsystem

```
root@ryzen01 ~ # ll /sys/devices/system/cpu/cpufreq/policy0
total 0
drwxr-xr-x  2 root root    0 Nov 11 17:57 .
drwxr-xr-x 15 root root    0 Nov 11 17:57 ..
-r--r--r--  1 root root 4096 Nov 11 17:57 affected_cpus
-r--r--r--  1 root root 4096 Nov 11 17:57 amd_pstate_highest_perf
-r--r--r--  1 root root 4096 Nov 11 17:57 amd_pstate_lowest_nonlinear_freq
-r--r--r--  1 root root 4096 Nov 11 17:57 amd_pstate_max_freq
-rw-r--r--  1 root root 4096 Nov 11 17:57 boost
-r--r--r--  1 root root 4096 Nov 11 17:57 cpuinfo_max_freq
-r--r--r--  1 root root 4096 Nov 11 17:57 cpuinfo_min_freq
-r--r--r--  1 root root 4096 Nov 11 17:57 cpuinfo_transition_latency
-r--r--r--  1 root root 4096 Nov 11 17:57 related_cpus
-r--r--r--  1 root root 4096 Nov 11 17:57 scaling_available_governors
-r--r--r--  1 root root 4096 Nov 11 17:57 scaling_cur_freq
-r--r--r--  1 root root 4096 Nov 11 17:57 scaling_driver
-rw-r--r--  1 root root 4096 Nov 11 17:57 scaling_governor
-rw-r--r--  1 root root 4096 Nov 11 17:57 scaling_max_freq
-rw-r--r--  1 root root 4096 Nov 11 17:57 scaling_min_freq
-rw-r--r--  1 root root 4096 Nov 11 17:57 scaling_setspeed
```

DVFS – P-State Management

Linux CPUFreq Subsystem

ondemand

```
(.menv) root@ryzen01 ~ # ll /sys/devices/system/cpu/cpufreq/ondemand/
total 0
drwxr-xr-x 2 root root 0 Nov 12 16:32 .
drwxr-xr-x 16 root root 0 Nov 11 17:57 ..
-rw-r--r-- 1 root root 4096 Nov 12 16:32 ignore_nice_load
-rw-r--r-- 1 root root 4096 Nov 12 16:32 io_is_busy
-rw-r--r-- 1 root root 4096 Nov 12 16:32 powersave_bias
-rw-r--r-- 1 root root 4096 Nov 12 16:32 sampling_down_factor
-rw-r--r-- 1 root root 4096 Nov 12 16:32 sampling_rate
-rw-r--r-- 1 root root 4096 Nov 12 16:32 up_threshold
(.menv) root@ryzen01 ~ #
```

schedutil

```
ewma_sum(u) := u_0 + u_1*y + u_2*y^2 + ...
ewma(u) = ewma_sum(u) / ewma_sum(1)

max( running, util_est ); if UTIL_EST
u_cfs := { running; otherwise

clamp( u_cfs + u_rt , u_min, u_max ); if UCLAMP_TASK
u_clamp := { u_cfs + u_rt; otherwise

u := u_clamp + u_irq + u_dl; [approx. see source for more detail]


f_des := min( f_max, 1.25 u * f_max )
```

- Takes current CPU load information from scheduler
- "Short-term" history has benefits and problems
- Typically executed every 10ms

- Exponential Weighted Moving Average (EWMA)
 $> 50\% \text{ last } 32\text{ms}, 50\% \text{ rest of history}$
- "Long-term" history less flexible
- Executed everytime with scheduler

DVFS – P-State Management

Transition Latency



Source: IBM PowerPC 405 LP NetSeminar

Marco Liess (TUM) | Linux Power Management | ecoCompute Conference 2025

DVFS – P-State Management

Power consumption of digital circuits

- 1) Dynamic power consumption:

$$P_{dyn} = \alpha C \cdot f_{clk} \cdot V_{dd}^2$$

α : activity, C: capacity,
 f_{clk} : clock frequency, V_{dd} : supply voltage

- 2) Available processing capacity:

$$MIPS_{prov} = N \cdot f_{clk} \cdot IPC$$

N: # of cores, f_{clk} : clock frequency, IPC: instructions per cycle

- 3) Required processing capacity

$$MIPS_{req} = ?$$

Depends on application and can be hard to predict!





Figure: Power consumption of one core of a Ryzen 5 5600G 6-core CPU

→ Use f_{clk} to match supply to demand!

Processing Demand in High-Speed Networking

The required processing capacity of NIC depends on:



Source: <https://iebmedia.com/wp-content/uploads/2023/05/Ethernet-Speeds.png>

Processing Demand in High-Speed Networking


The required processing capacity of NIC depends on:



Packet rate / Packet interarrival time

$$\rightarrow \text{packets per second (PPS)} \quad PPS = \frac{\text{link_speed}}{\text{packet_size}}$$

Link speed	40 Gbps		100 Gbps		800 Gbps	
Size [Byte]	64	512	64	512	64	512
PPS [M]	78	9.8	195	24	1,563	195
1 / PPS [ns]	12.8	102	5.1	41	0.64	5.1




Source: <https://iebmedia.com/wp-content/uploads/2023/05/Ethernet-Speeds.png>

$$\text{Date_rate} = w_{\text{Data}} \times f$$


Processing Demand in High-Speed Networking

Header Processing




- Fixed header location allows simple parsing
- **Common use-cases:** routing, switching, IP/port-based firewalls, ...

Payload Processing



- **Common use-cases:** application/session/user identification for firewalls or bandwidth throttling, cryptology, intrusion detection, virus scanning, ...



Source: Netronome


Processing Demand in High-Speed Networking



The processing complexity of a networking function:

→ *instructions per packet (IPP)*

$$\text{Required processing capacity } \text{IPS} = \text{PPS} * \text{IPP} \left[\frac{\text{instructions}}{\text{second}} \right]$$



Source: Netronome

@ 100Gbps

$$\text{L2 Switching (64B): } \text{PPS} * \text{IPP} = 195.313.000 \frac{\text{packets}}{\text{second}} * 75 \frac{\text{instructions}}{\text{packet}} = 14.6 * 10^9 \frac{\text{instructions}}{\text{second}} (\text{GIPS})$$

$$\text{Intrusion Detection (512B): } 146 * 10^9 \frac{\text{instructions}}{\text{second}} \quad \text{IPSec (512B) : } 403 * 10^9 \frac{\text{instructions}}{\text{second}} (\text{GIPS})$$

Processing Demand in High-Speed Networking



AMD Ryzen 5 5600: 6 cores @ 3.9 GHz, 65 W TDP

$$IPS = 6 \text{ cores} * f_{clk} * IPC_{core} \approx 23,4 * 10^9 \frac{\text{instructions}}{\text{second}}$$

$$IPC_{core} \approx 1$$

L2 Switching ✓

Intrusion Detection ✗

IPSec ✗

Networking Function	100Gbps
L2 Switching	$14.6 * 10^9$
Intrusion Detection	$146 * 10^9$
IPSec	$403 * 10^9$

IPSec @ 100Gbps would require 18 CPUs with a TDP of 1.1 kW!



DVFS – P-State Management

Power consumption of digital circuits

- 1) Dynamic power consumption:

$$P_{dyn} = \alpha C \cdot f_{clk} \cdot V_{dd}^2$$

α : activity, C: capacity,
 f_{clk} : clock frequency, V_{dd} : supply voltage

- 2) Available processing capacity:


$$MIPS_{prov} = N \cdot f_{clk} \cdot IPC$$

N : # of cores, f_{clk} : clock frequency, IPC : instructions per cycle

- 3) Required processing capacity:

$$MIPS_{req} = MPPS \cdot IPP$$

$MPPS$: million packets per second, IPP : instructions per packet



Oeldemann, A., Wild, T., & Herkersdorf, A. (2017, March). Reducing data center resource over-provisioning through dynamic load management for virtualized network functions. In *International Conference on Architecture of Computing Systems* (pp. 234-247). Cham: Springer International Publishing.

→ Use f_{clk} to match supply to demand!

DVFS – P-State Management

Power consumption of digital circuits

- 1) Dynamic power consumption:

$$P_{dyn} = \alpha C \cdot f_{clk} \cdot V_{dd}^2$$

α : activity, C: capacity,
 f_{clk} : clock frequency, V_{dd} : supply voltage

- 2) Available processing capacity:


$$MIPS_{prov} = N \cdot f_{clk} \cdot IPC$$

N: # of cores, f_{clk} : clock frequency, IPC: instructions per cycle

- 3) Required processing capacity:

$$MIPS_{req} = MPPS \cdot IPP$$


MPPS: million packets per second, IPP: instructions per packet



→ Use f_{clk} to match supply to demand!


Analytic Model

Timeline visualization of DVFS




Analytic Model

Timeline visualization of DVFS




Analytic Model

Timeline visualization of DVFS



Analytic Model

Learnings



- Traffic fluctuations and flow diversity lead to quickly changing processing demand
- Event-driven, short-term compute demand difficult to manage for OS
- Minimal reaction time necessary to
 - Achieve maximum energy savings
 - Minimize negative impact on latency

Liess, M., Demicoli, J., Tiedje, T., Lohrmann, M., Nickel, M., Luniak, M., ... & Herkersdorf, A. (2023, November). X-mape: Extending 6g-connected self-adaptive systems with reflexive actions. In *2023 IEEE Conference on Network Function Virtualization and Software Defined Networks (NFV-SDN)* (pp. 163-167). IEEE.

SmartNICs

Motivation



Performance and Energy Efficiency

- cope with very high data rates (up to hundreds of Gbps)
- lowest packet delay as possible
- low power consumption



Customized **ASIC**




Flexibility

- adapt to evolving packet processing applications
- efficient resource sharing among network applications



Programmable **CPU / ASIP**



Source: T. Noll / H. Blume et al., „Model-based exploration of the Design Space for Heterogeneous System on Chip“, 2002

SmartNICs

Motivation



Performance and Energy Efficiency

- cope with very high data rates (up to hundreds of Gbps)
- lowest packet delay as possible
- low power consumption



Customized **ASIC**




Flexibility

- adapt to evolving packet processing applications
- efficient resource sharing among network applications




Programmable **CPU / ASIP**



Source: T. Noll / H. Blume et al., „Model-based exploration of the Design Space for Heterogeneous System on Chip“, 2002

SmartNIC Extensions @ LIS




Linux Network Stack


OS Support?

Challenges:

- Driver, Network Stack and App not „synchronized“
- Buffers disaggregate between functions
→ Limited info about state of other components
- Processing demand highly dependent on App
→ Difficult to predict without feedback



NETTB - NETworking TestBed @ LIS




NETTB - NETworking TestBed @ LIS


1. Configuration

```
! benchmark_settings.yaml > gen_trace
1   dut_pc: ryzen01
2   dut_benchmark: l2fwd
3   dut_threads: [6,12]
4   dut_pstate: [ondemand,performance]
5   gen_pc: ryzen02
6   gen_trace: caida
```


3. Results



2. Running




4. Postprocessing



Experimental Evaluation


Experimental Setup

CAIDA Network Trace
2019 (10Gbps)





Experimental Evaluation

Linux Governors






Experimental Evaluation

Hardware Priority Pinning




Experimental Evaluation

Hardware Priority Pinning




Experimental Evaluation

Linux Governors



Experimental Evaluation



Linux Governors



Experimental Evaluation

ecoNIC: SmartNIC-assisted Power Management ¹

Combines traffic-dependent power management in OS / Linux
with priority-aware traffic steering / pinning in SmartNIC HW





Liess, M., Biersack, F., Nolte, L., Wild, T., & Herkersdorf, A. (2025). ecoNIC: SmartNIC-assisted power management for networking workloads in Linux servers. *Micropocessors and Microsystems*, 105209.

Experimental Evaluation

ecoNIC: SmartNIC-assisted Power Management¹

Combines traffic-dependent power management in OS / Linux
with priority-aware traffic steering / pinning in SmartNIC HW




Liess, M., Biersack, F., Nolte, L., Wild, T., & Herkersdorf, A. (2025). ecoNIC: SmartNIC-assisted power management for networking workloads in Linux servers. *Micropocessors and Microsystems*, 105209.

Experimental Evaluation

ecoNIC: SmartNIC-assisted Power Management¹


Combines traffic-dependent power management in OS / Linux with priority-aware traffic steering / pinning in SmartNIC HW



Liess, M., Biersack, F., Nolte, L., Wild, T., & Herkersdorf, A. (2025). ecoNIC: SmartNIC-assisted power management for networking workloads in Linux servers. *Micropocessors and Microsystems*, 105209.

Experimental Evaluation


ecoNIC: SmartNIC-assisted Power Management ¹



Liess, M., Biersack, F., Nolte, L., Wild, T., & Herkersdorf, A. (2025). ecoNIC: SmartNIC-assisted power management for networking workloads in Linux servers. *Micropocessors and Microsystems*, 105209.

Experimental Evaluation


ecoNIC: SmartNIC-assisted Power Management ¹



Liess, M., Biersack, F., Nolte, L., Wild, T., & Herkersdorf, A. (2025). ecoNIC: SmartNIC-assisted power management for networking workloads in Linux servers. *Micropocessors and Microsystems*, 105209.

Experimental Evaluation

ecoNIC: SmartNIC-assisted Power Management ¹



Liess, M., Biersack, F., Nolte, L., Wild, T., & Herkersdorf, A. (2025). ecoNIC: SmartNIC-assisted power management for networking workloads in Linux servers. *Microprocessors and Microsystems*, 105209.

Conclusions & Outlook

- Power Management essential for energy-efficient server operation
 - Fluctuating, unpredictable network processing demand poses huge challenge
 - CPU / OS power management not optimized for external demand
- SmartNICs can monitor traffic demand right at the source
→ offers improvement in adaptability and energy-efficiency

Outlook:

- Extend SmartNIC with workload prediction, application information, etc.



Let's discuss!

Acknowledgement

We are greatful for the financial support by the *Bavarian Ministry of Economic Affairs, Regional Development and Energy* in the *6G Future Lab Bavaria* project and the *Federal Ministry of Education and Research* of Germany in the *6G-life* project with project identification number: 16KISK002

Bayerisches Staatsministerium für
Wirtschaft, Landesentwicklung und Energie



GEFÖRDERT VOM



Bundesministerium
für Bildung
und Forschung