



UNIVERSITÄT
LEIPZIG

Green Configuration

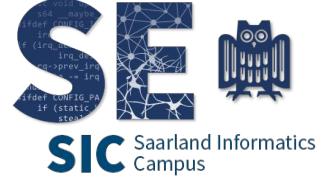
Energy Consumption of Configurable Software Systems



Max Weber



UNIVERSITÄT
LEIPZIG



Alina Mailach



Johannes Dorn



Norbert Siegmund



Max Weber



Florian Sattler



Sven Apel



Christian Kaltenecker



Deutsche
Forschungsgemeinschaft

Projekt goal: Assessing and reducing energy
consumption of configurable software systems

The climate is warming



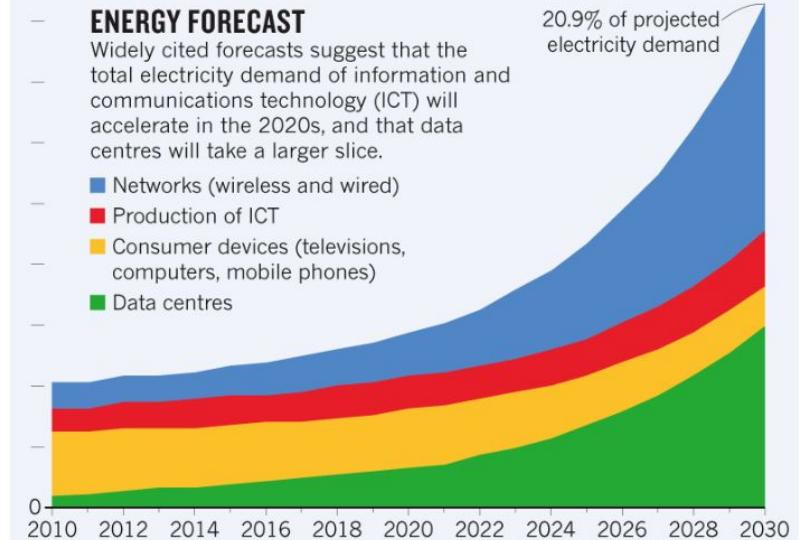
Energy consumption is rising

9,000 terawatt hours (TWh)

ENERGY FORECAST

Widely cited forecasts suggest that the total electricity demand of information and communications technology (ICT) will accelerate in the 2020s, and that data centres will take a larger slice.

- Networks (wireless and wired)
- Production of ICT
- Consumer devices (televisions, computers, mobile phones)
- Data centres



Energy Consumption at ecoCompute



Workloads, Use Cases

Software Systems

OS, Languages

Hardware



Rerouting the Cloud:
Cutting 90% CO₂ and
Cloud Costs with Smarter
Workload Shifting

Dryden Williams
CarbonRunner



Coding Smarter for a
Greener Future | A
comparison between Go
& Java

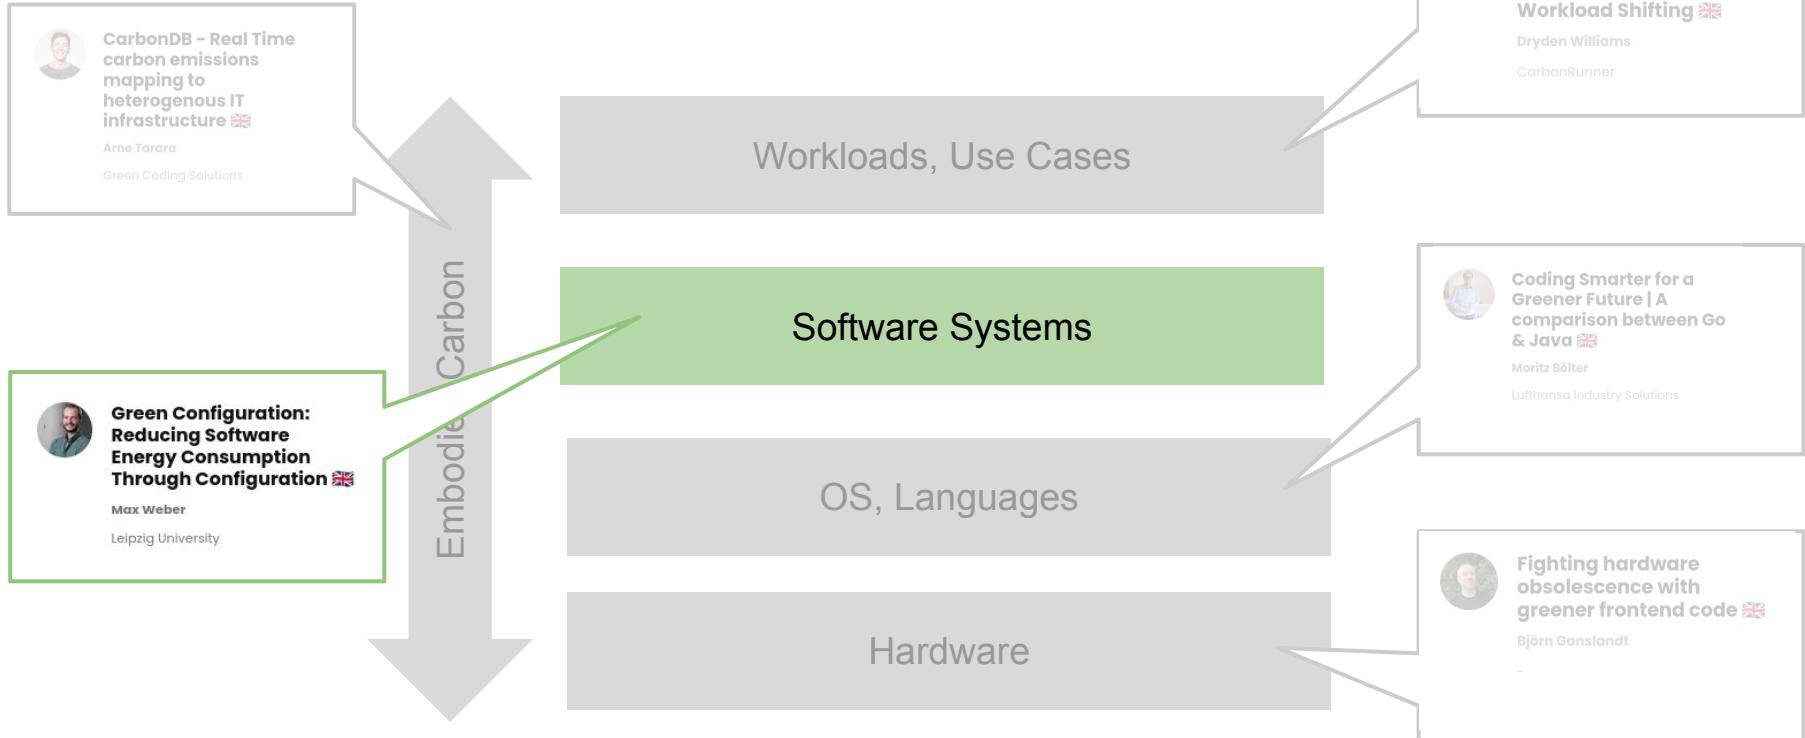
Moritz Böller
Lufthansa Industry Solutions



Fighting hardware
obsolescence with
greener frontend code

Björn Ganslandt

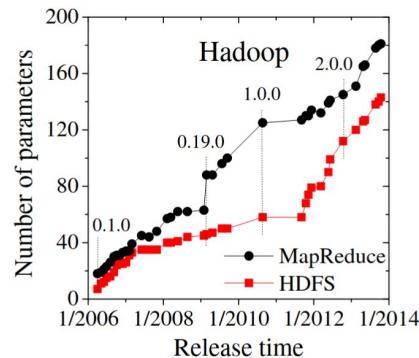
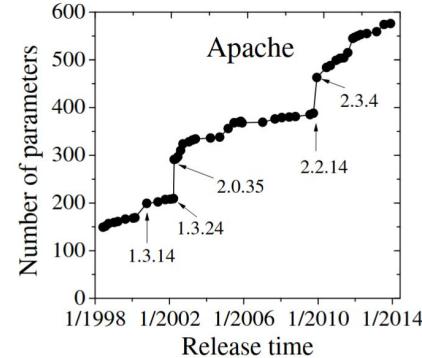
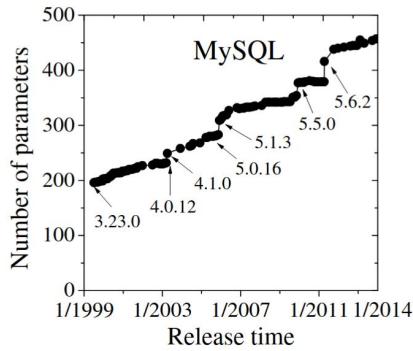
Energy Consumption at ecoCompute



How can we reduce the energy consumption of software systems?

By tailoring software to its workload and usage scenario.

Configuration is everywhere!



Today's available configuration decisions:

MySQL – 700+

Apache – 900+

Hadoop – 400+

Most users stick to default configurations,
leaving much potential unused.

1

Assessing Energy Consumption



How can we **measure** software energy consumption?

2

Debugging Energy consumption



3

Practitioners' perspective



Energy Consumption of Configurable Software Systems

Measuring Energy Consumption at the Code Level

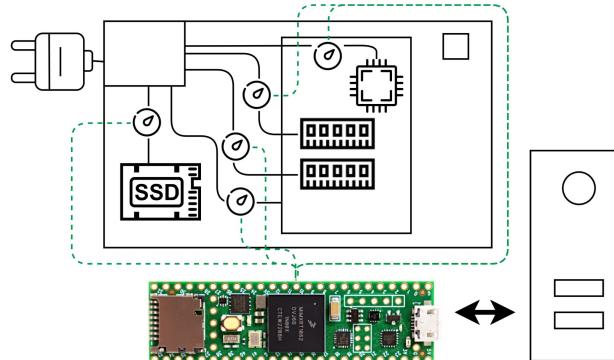


component-wise measurements
(adding up to the whole system)

High sampling rate:
3.5 kHz

no energy measurement bias

Accuracy: < 1.2%
measurement error



Hardware Components

Raspberry Pi 4B



60€

USB

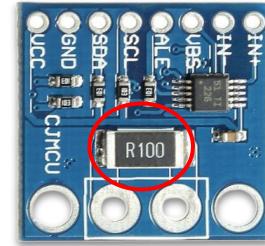
Teensy 4.1



30€

I²C

INA 226
Sample Time 280 us
Sample Rate 3.5 kHz



5€

3,500 measurements per second
≈ 1 sample every 0.3 ms (for all components)

175€

Hardware Components

Raspberry Pi 4B



60€

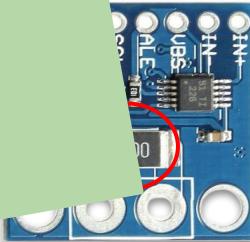
USB

Teensy 4.1



Hours it took until it worked reliably:
Priceless.

INA 226
Sample Time 280 us
Sample Rate 3.5 kHz

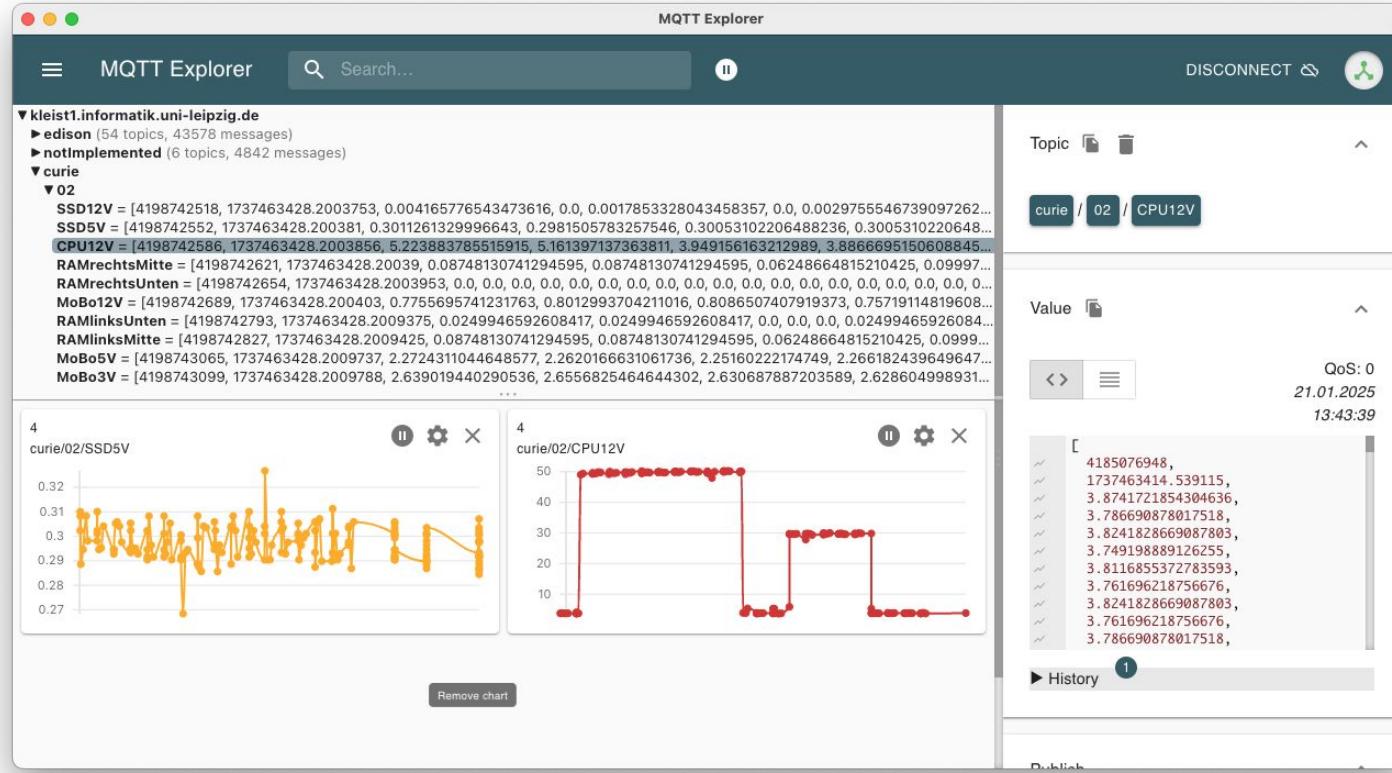


5€

3,500 measurements per second
 $\hat{=} 1$ sample every 0.3 ms (for all components)

175€

Power Measurement

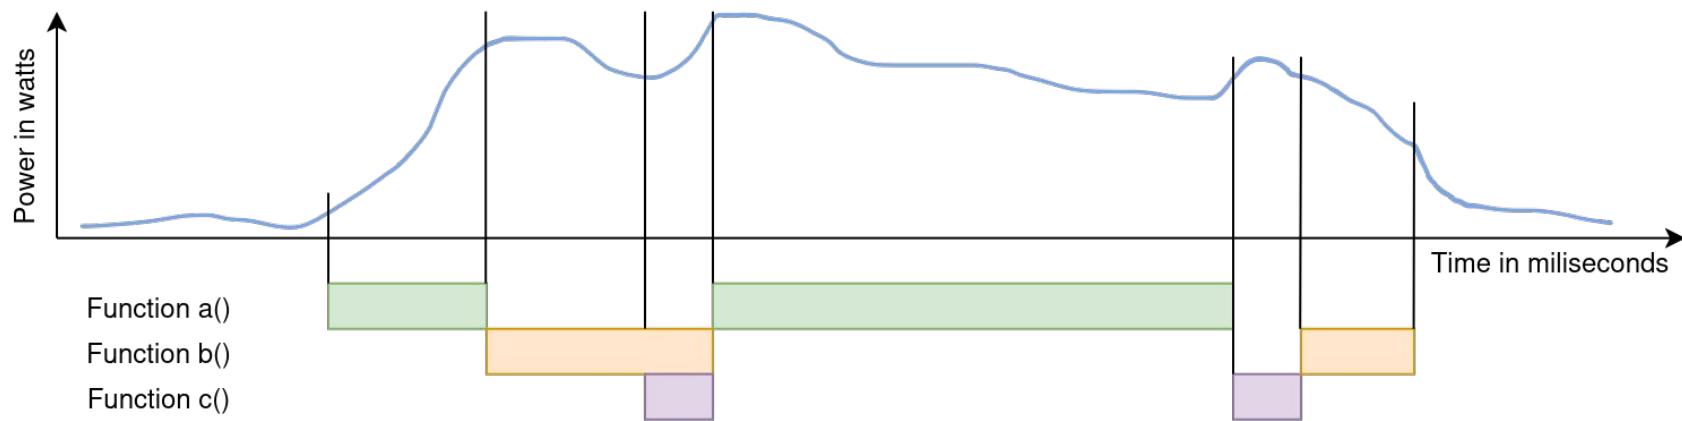


Performance Profiling

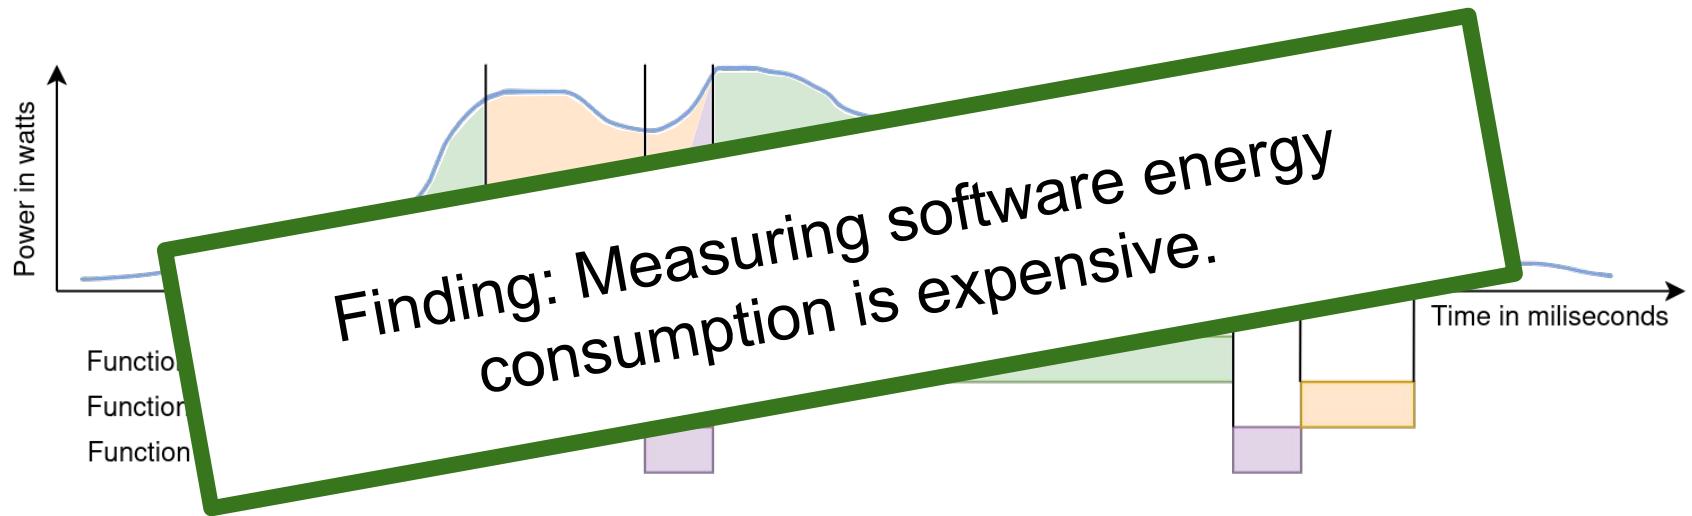
Linux *perf*



Synchronizing Energy and Performance Measurements



Synchronizing Energy and Performance Measurements



1

Assessing Energy Consumption



How can we **measure** software energy consumption?

Can we find a **reliable proxy**?

Energy Consumption of Configurable Software Systems

2

Debugging Energy consumption



3

Practitioners' perspective



Energy Metering is expensive, time consuming
and tool intensive.



Infeasible in many projects

Solution: Using performance as a proxy for energy consumption



System-Level Correlation

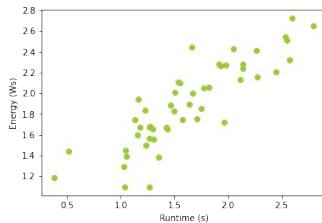


Option-Level Correlation



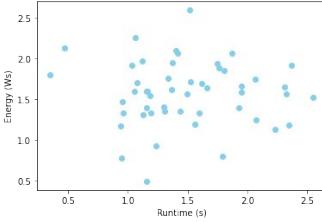
Literature Study Statements

positive correlation



“perform significantly better and consume less energy with only a small loss in QoS” [1]

no / absent correlation

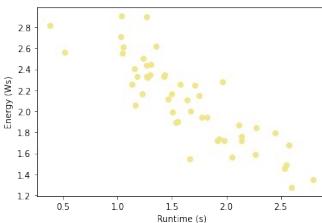


“the involved physical components may not consume energy proportional to time” [2]

“caching has an effect on performance but not on energy consumption” [2]

“no correlation if some communication comes into play” [4]

negative correlation



“computational offloading brings better performance but not always better energy efficiency” [3]

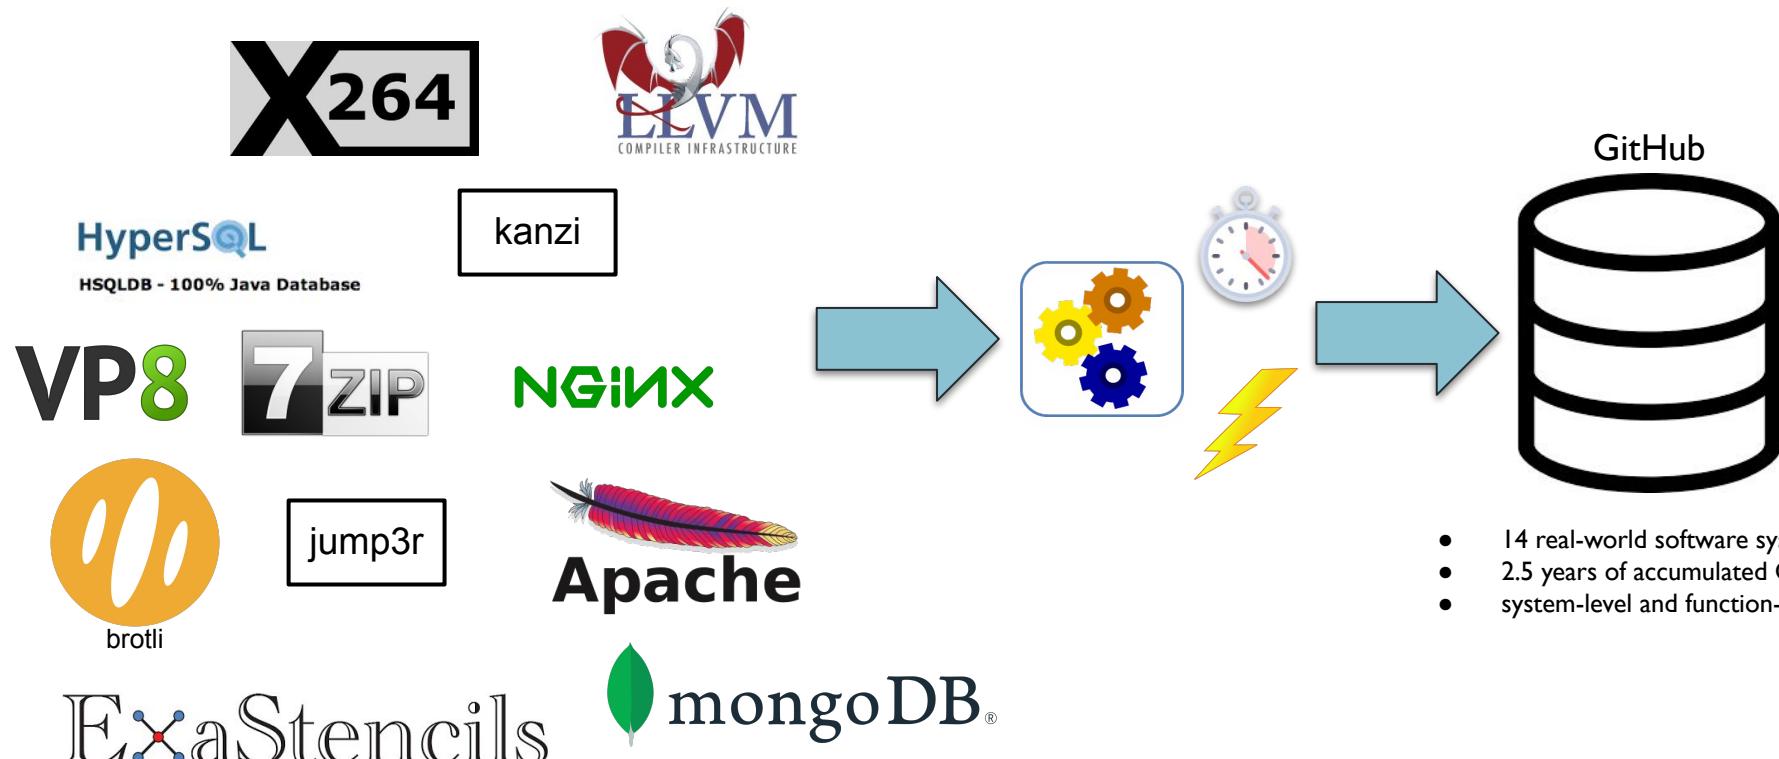
[1] Green: a framework for supporting energy-conscious programming using controlled approximation, 2020

[2] Evaluating the Impact of Caching on the Energy Consumption and Performance of Progressive Web Apps, 2020

[3] Analysis of Performance and Energy Consumption of Wearable Devices and Mobile Gateways in IoT Applications, 2019

[4] Quantifying energy use in dense shared memory HPC node, 2016

Empirical Study to Investigate Performance-Energy Correlation



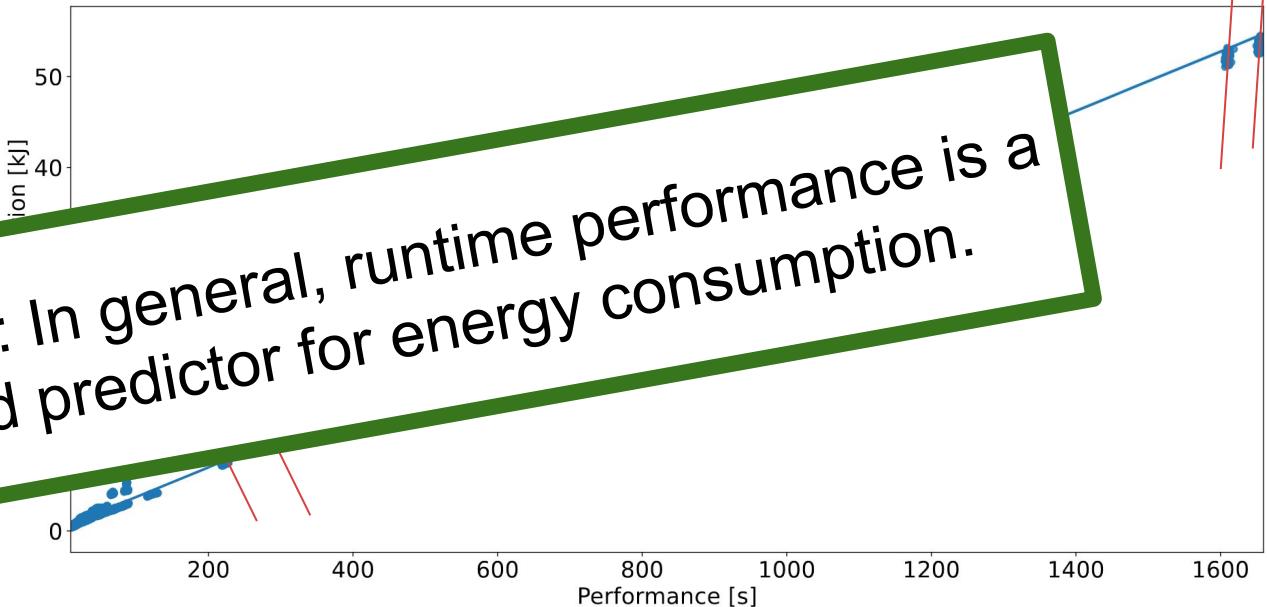
System-Level Correlation



Correl.

lzip

Finding: In general, runtime performance is a good predictor for energy consumption.



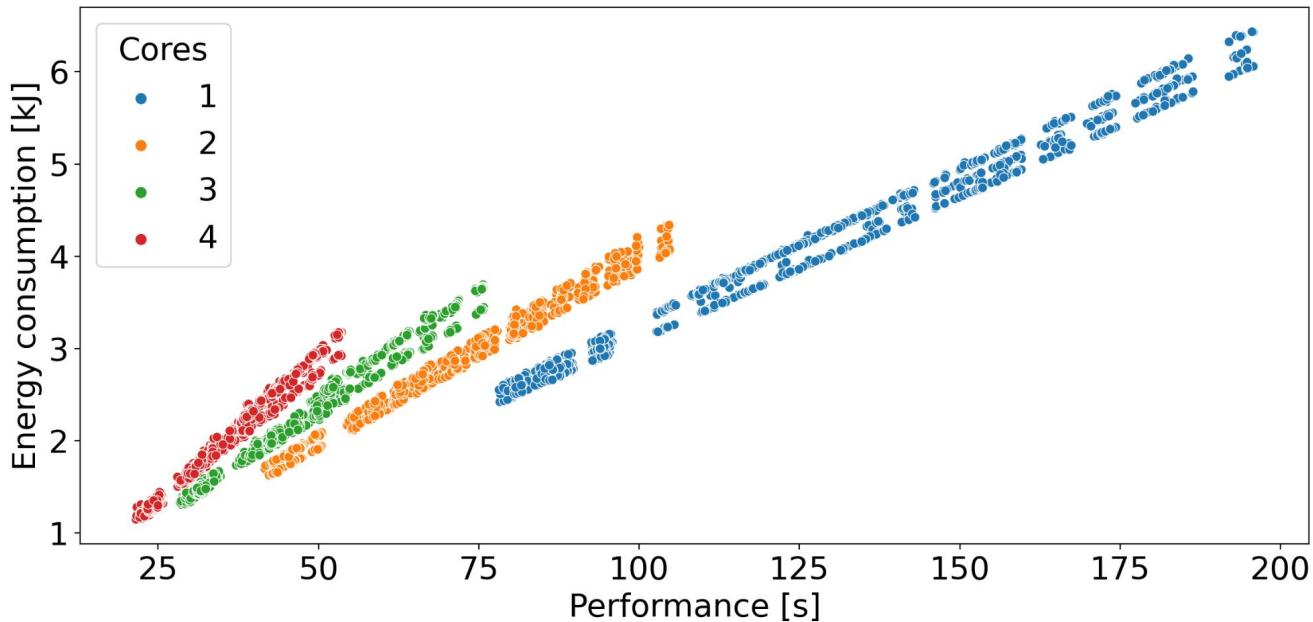
Option-Level Correlation



Correlation: 0.99
MAPE: 10.3%



Correlation: 0.99
MAPE: 2.5%



Option-Level Correlation

HyperSQL

HSQLDB - 100% Java Database



Correlation:
MAPE:

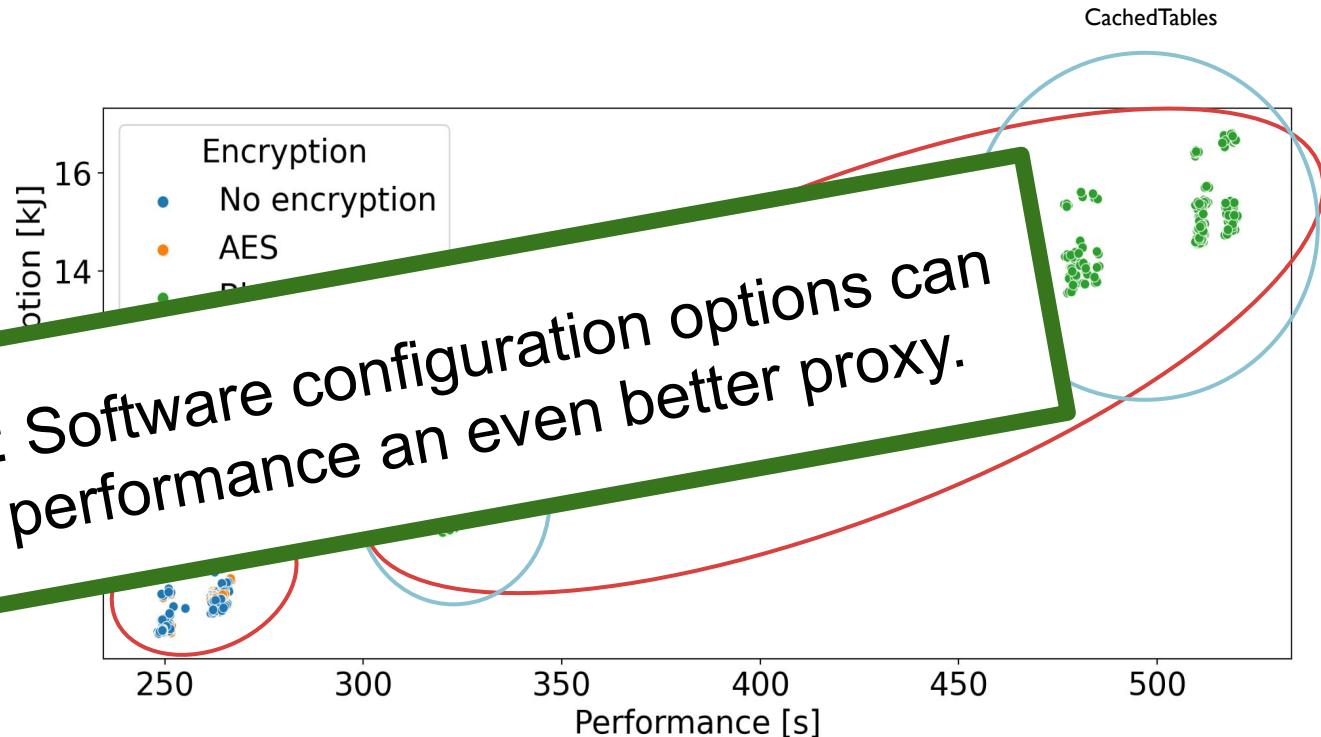


Finding: Software configuration options can
make performance an even better proxy.

Blowfish & Memory Tables



Correlation: 0.13
MAPE: 213.5%



1

Assessing Energy Consumption



How can we **measure** software energy consumption?

Can we find a **reliable proxy**?

How can we get **code-level insights**?

Energy Consumption of Configurable Software Systems

2

Debugging Energy consumption

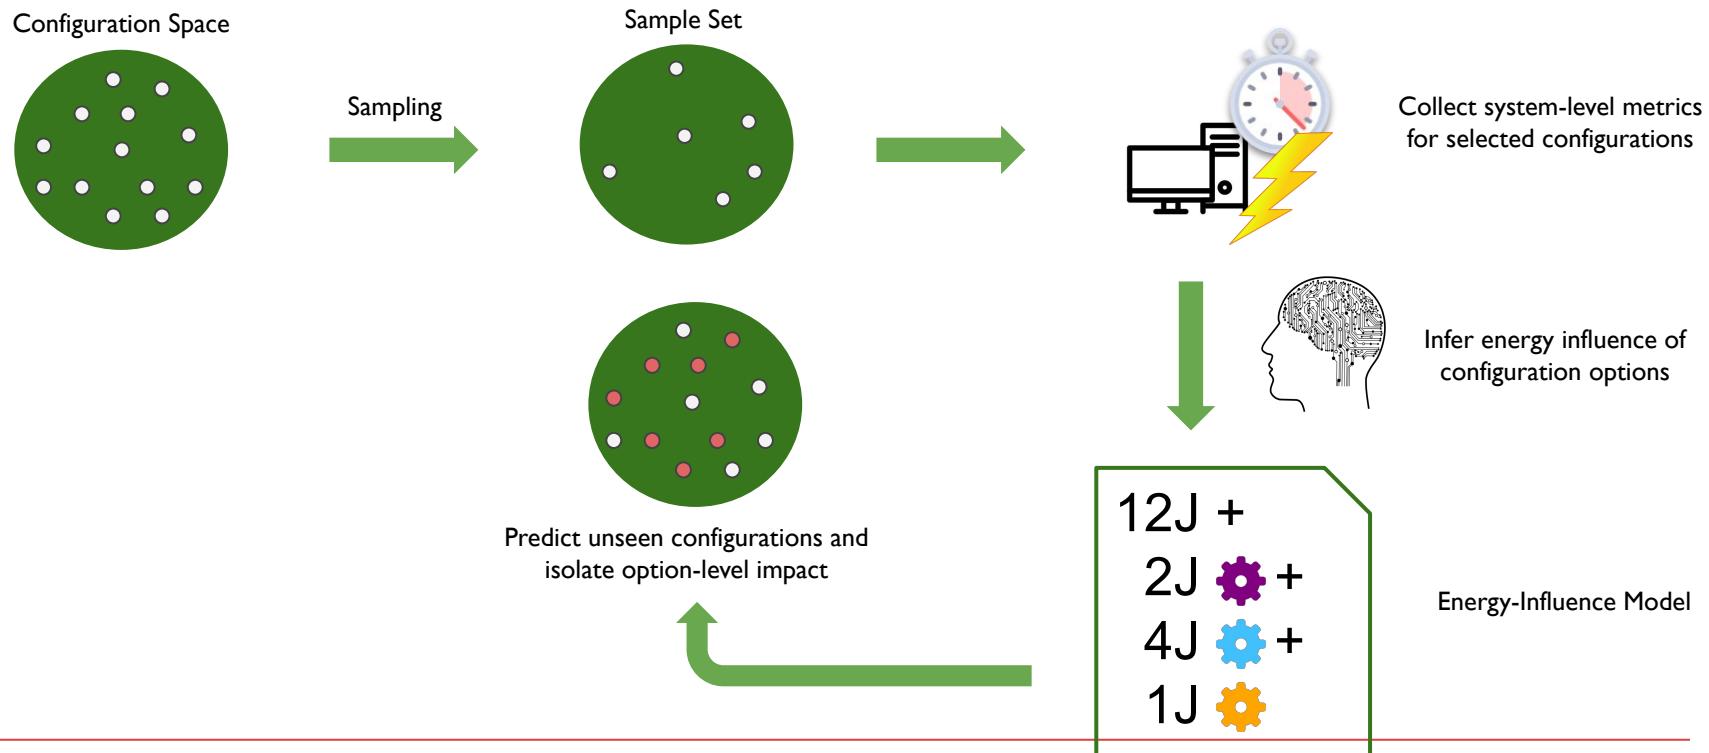


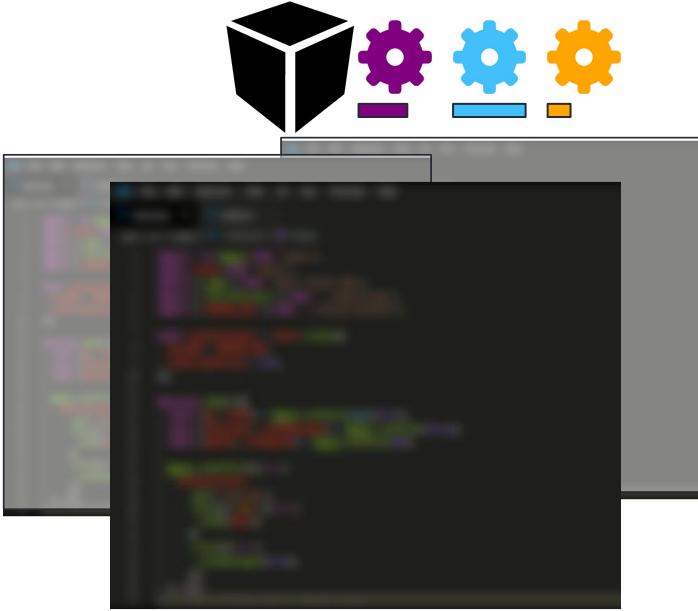
3

Practitioners' perspective



Building Energy-Influence Models for Configurable Systems



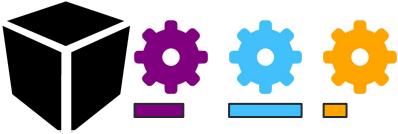


As a **developer**, how do I know which part of my code I need to change?

Black-Box Models

As a **user**, I can optimize configurations to reduce energy consumption 💪

... but I can't see **where** in the code the energy goes.



How can we **efficiently** pin point energy hotspots in the code base with minimal measurement effort?

Black-Box Models

As a **user**, I can optimize configurations to reduce energy consumption 💪

... but I can't see **where** in the code the energy goes.

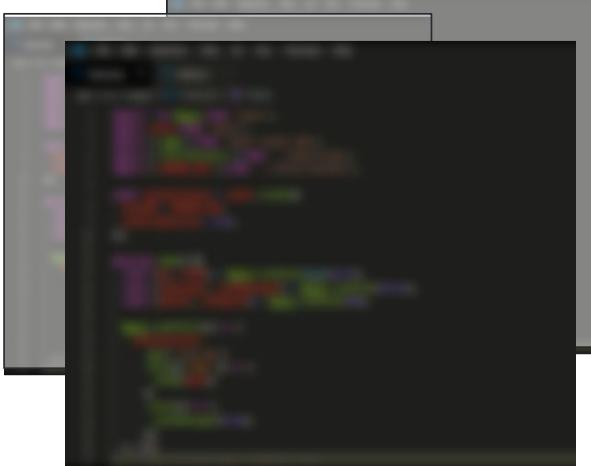
A screenshot of a code editor showing a file named 'home.tsx'. The code is a React component that makes an axios request to '/user/me' and updates state based on the response. Line 18, which contains the axios call, is highlighted with a red box.

```
1 import * as React from 'react';
2 import axios from 'axios';
3 import { Link } from 'react-router-dom';
4 import { startCheckout } from './lib/stripe';
5 import { SERVER_URL } from '../cils/contents';
6
7 const axiosInstance = axios.create({
8   baseURL: SERVER_URL,
9   withCredentials: true,
10 });
11
12 function Home() {
13   const [me, setMe] = React.useState<any>(null);
14   const [showLogin, setShowLogin] = React.useState(false);
15   const [amount, setAmount] = React.useState(10);
16
17   React.useEffect(() => {
18     axiosInstance
19       .get('/user/me')
20       .then(({ data }) => {
21         setMe(data);
22       })
23       .catch(() => {
24         setShowLogin(true);
25       });
26   }, []);
}
```

Established White-Box Models

We know exactly **which** code statements consume energy.

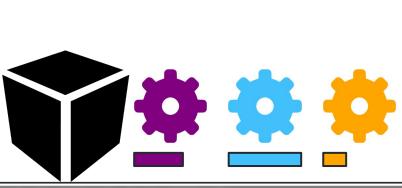
...but measuring this across a full system is often **infeasible** 😞



Black-Box Models

As a **user**, I can optimize configurations to reduce energy consumption 💪

... but I can't see **where** in the code the energy goes. 😞



```

1 import * as React from 'react';
2 import axios from 'axios';
3 import { Link } from 'react-router-dom';
4 import { startCheckout } from './lib/stripe';
5 import { SERVER_URL } from './utils/constants';
6
7 const axiosInstance = axios.create({
8   baseURL: SERVER_URL,
9   withCredentials: true,
10 });
11
12 function Home() {
13   const [me, setMe] = React.useState<any>(null);
14   const [showLogin, setShowLogin] = React.useState(false);
15   const [amount, setAmount] = React.useState(10);
16
17   React.useEffect(() => {
18     axiosInstance
19       .get('/user/me')
20       .then(({ data }) => {
21         setMe(data);
22       })
23       .catch(() => {
24         setShowLogin(true);
25       });
26   }, []);

```

Method-Level White-Box Models

As a user, I still benefit from efficient configurations that save energy ⚡.

As a developer, I can identify which methods consume energy without measuring every single line of code.

```

1 import * as React from 'react';
2 import axios from 'axios';
3 import { Link } from 'react-router-dom';
4 import { startCheckout } from './lib/stripe';
5 import { SERVER_URL } from './utils/constants';
6
7 const axiosInstance = axios.create({
8   baseURL: SERVER_URL,
9   withCredentials: true,
10 });
11
12 function Home() {
13   const [me, setMe] = React.useState<any>(null);
14   const [showLogin, setShowLogin] = React.useState(false);
15   const [amount, setAmount] = React.useState(10);
16
17   React.useEffect(() => {
18     // ...
19     .get('/user/me')
20     .then(({ data }) => {
21       setMe(data);
22     })
23     .catch(() => {
24       setShowLogin(true);
25     });
26   }, []);

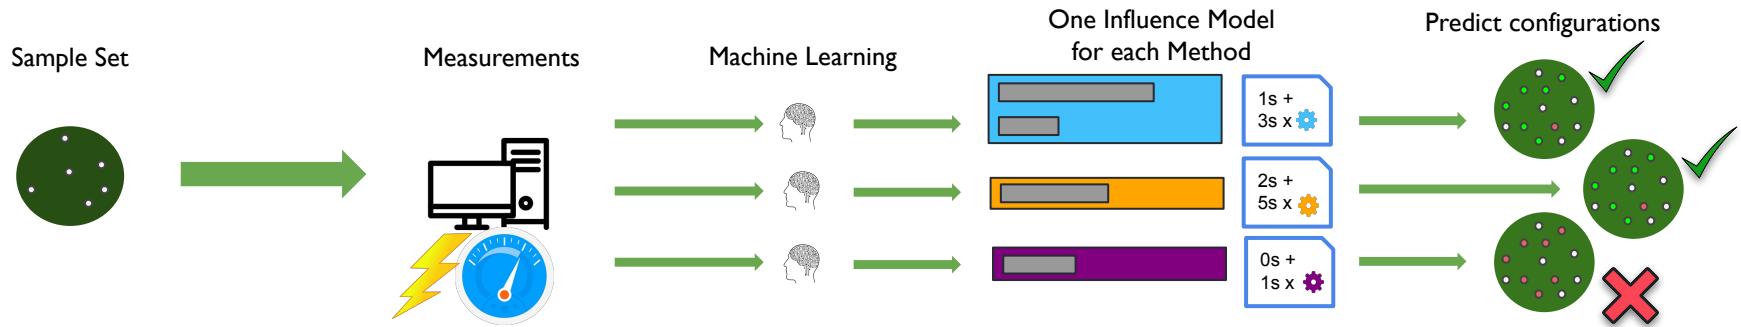
```

Established White-Box Models

We know exactly **which** code statements consume energy. 💪

...but measuring this across a full system is often **infeasible**. 😞

White-Box Energy Influence Models via Profiling



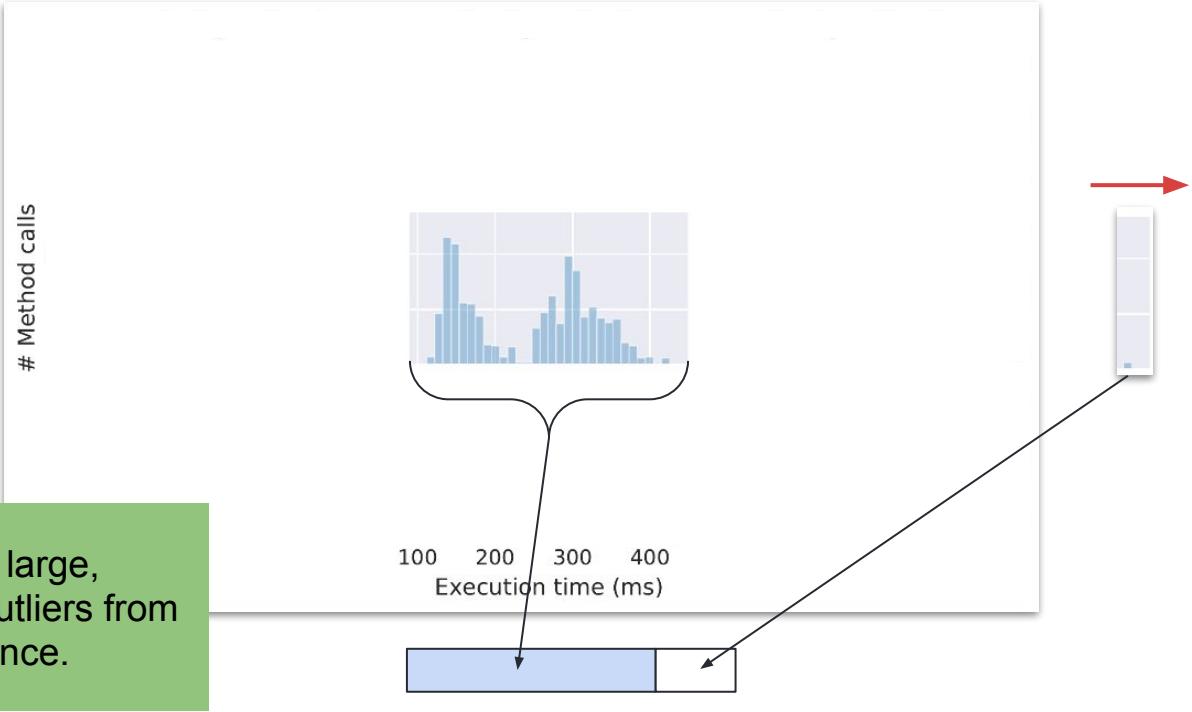
Challenge: Handling Variance in Measurement Data

Context variance

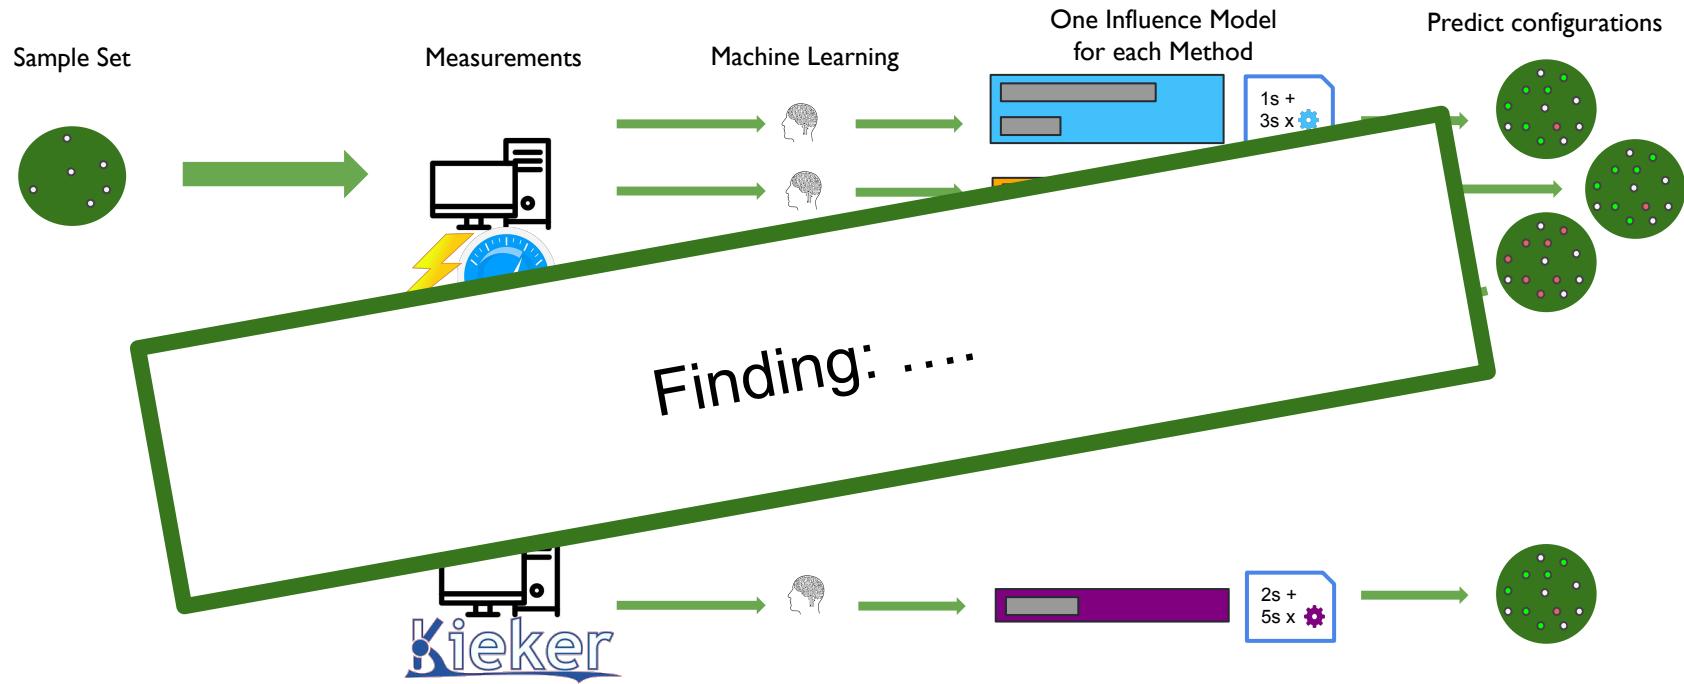
Configuration variance

Measurement variance

Solution: Filter large,
non-deterministic outliers from
context variance.



White-Box Energy Influence Models via Profiling



IDE Integration

The screenshot shows the Eclipse Platform interface with the title "runtime-EclipseApplication - catena/src/main/java/components/graph/algorithms/GenericGraph.java - Eclipse Platform".

Package Explorer: Shows the project structure under "catena [GreenDev master]". The file "DoubleButterflyGraph.java" is selected.

Code Editor: Displays the Java code for "GenericGraph.java". The code implements a graph algorithm using byte arrays and helper objects.

```
public byte[][] graph (int g, byte[][] v, int lambda){  
    int dim1 = (int) Math.pow(2, g);  
    int dim2 = hPrime.getOutputSize();  
    byte[][] r = new byte[dim1][dim2];  
  
    for (int k = 0; k < lambda; ++k){  
        r[0] = hFirst(helper.concatenateByteArrays(v[dim1-1],  
            v[indexing.getIndex(0, g)]));  
  
        int loop = (int) Math.pow(2, g);  
  
        for (int i = 1; i < loop; ++i){  
            hPrime.update(helper.concatenateByteArrays(r[i-1], v[indexing.getIndex(i, g)]));  
            r[i] = hPrime.doFinal();  
        }  
  
        System.arraycopy(r, 0, v, 0, r.length);  
    }  
    return v;  
}  
  
private byte[] hFirst(byte[] in){  
    int n = h_.getOutputSize();  
    int k = hPrime.getOutputSize();  
    int l = k/n;  
  
    byte[][] w = new byte[l][n];  
    byte[] iByte = new byte[l];  
}
```

Performance Overview: A table showing performance values and paths.

Performance Value	Path
100.0	main.java.components.graph.algorithms.GenericGraph:graph
1.061113	main.java.Catena:flap
0.19015925	profiling.CatenaProfiler:initCatenaByConfig
0.1401815	main.java.components.gamma.algorithms.SaltMix:gamma

1

Assessing Energy Consumption



How can we **measure** software energy consumption?

Can we find a **reliable proxy**?

How can we get **code-level insights**?

Energy Consumption of Configurable Software Systems

2

Debugging Energy consumption



How can we identify **energy bugs**?



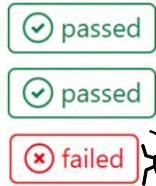
3

Practitioners' perspective



Performance Bugs in Configurable Systems

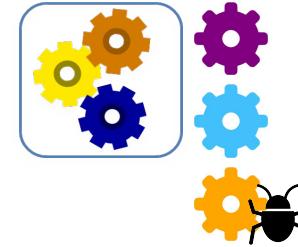
Functional bug



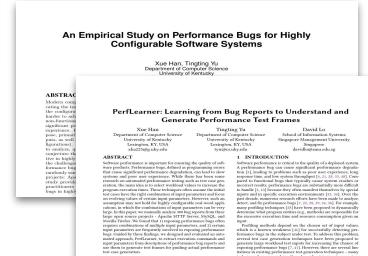
Performance bugs



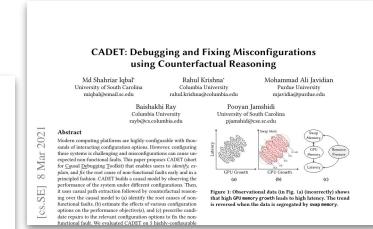
Configuration-dependent bugs



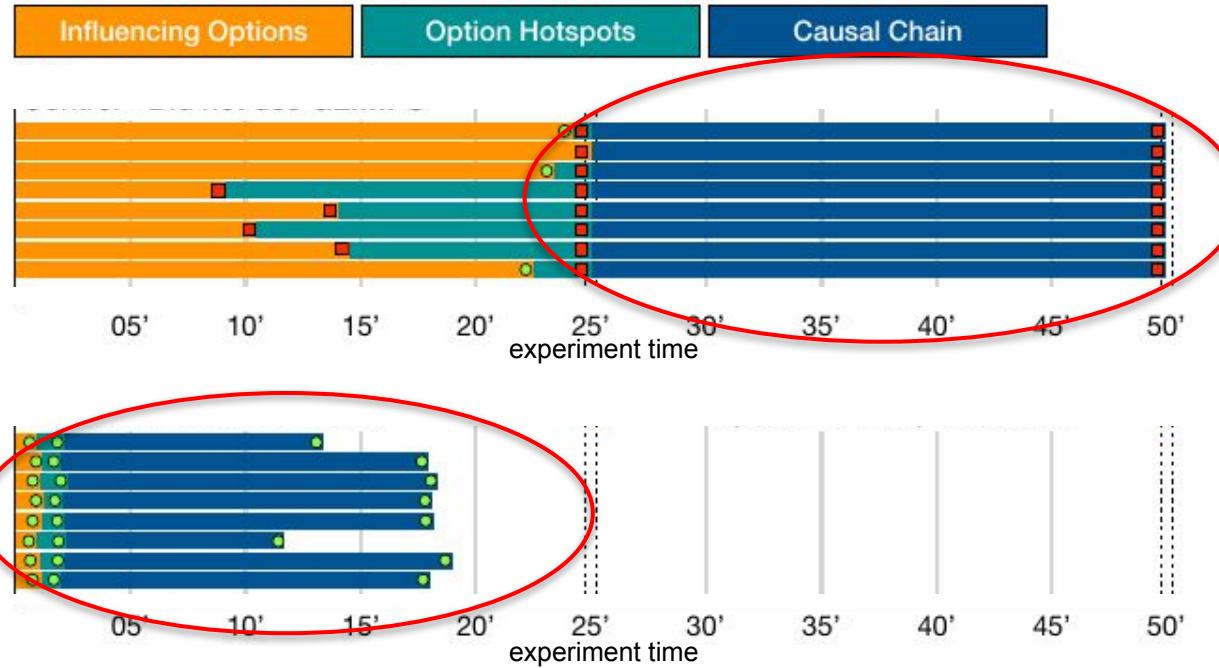
Requires fixing code



Fixes are time consuming



Why Debugging Performance Bugs Takes So Much Time



On Debugging the Performance of Configurable Software Systems: Developer Needs and Tailored Tool Support

Miguel Velez, Pooyan Jamshidi, Norbert Siegmund, Sven Apel, Christian Kästner

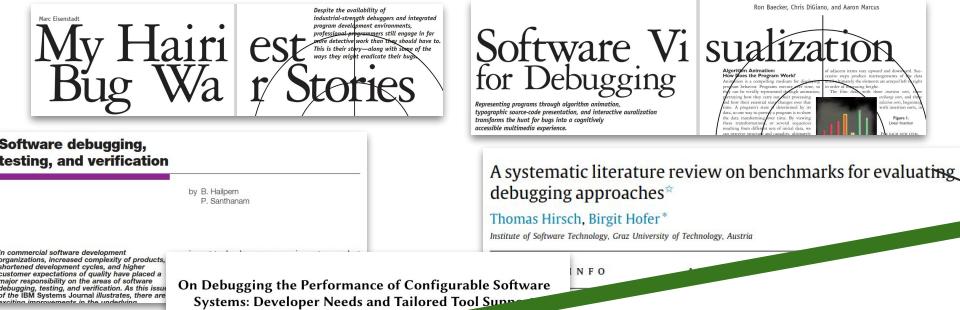
Carnegie Mellon University, University of South Carolina, Leipzig University, Stanford University, Carnegie Mellon University

ABSTRACT
Determining whether a configurable software system has a performance bug or not is often challenging. While there are numerous debugging techniques that can support developers in this process, they are often ineffective. One way to address the actual needs that developers have when debugging the performance of their systems is to support them with tailored tools. In this paper, we take a human-centered approach to understand developer needs and requirements for tool support developers in the process of debugging the performance of configurable software systems. We conducted a case study with 19 developers to identify the information needs that developers have when debugging the performance of their systems. We found that developers need a tool that implements a tailored tool, adapting techniques from prior work, to support these needs. Two user studies, with a total of 16 developers, evaluated the usefulness of the proposed tool. The results show that the tool is useful for developers to support them when debugging the performance of configurable software systems.

V1 (cc-SE) 19 Mar 2022

[Miguel Velez, Pooyan Jamshidi, Norbert Siegmund, Sven Apel, and Christian Kästner. On debugging the performance of configurable software systems: Developer needs and tailored tool support]

Debugging Strategies in the Literature



overview of debugging strategies

Title	Description	Actions	References
Program comprehension	Developers gain comprehensive and high-level knowledge of the system by browsing the code and files, using available hot-spot information.	hot-spot examination	[Murphy et al. 2008; Romero et al. 2007; Velez et al. 2022; Vessey 1985]
Inspect source code	Examining the source code to identify errors and bugs.	Examination	[Böhme et al. 2017; Hirsch and Hofer 2021; Murphy et al. 2008; Peng et al. 2016; Subrahmanian et al. 2008]
Follow data flow	Tracing the flow of data through the system to understand how it transforms within the system, thereby identifying potential issues.	follow data flow	[Subrahmanian et al. 2008]
Follow control flow	Tracking the program's execution path to understand the sequence of operations and identify deviations from expected behavior.	follow control flow	[Eisenstadt 1997; Gould 1975; Perscheid et al. 2017; Romero et al. 2007]
Follow abstraction	Transforming the system into different levels of abstraction to understand how it works at various levels.	follow abstraction	[Böhme et al. 2017; Hirsch and Hofer 2021; Murphy et al. 2008]

Finding: Literature is not consolidated!

- ❖ reviewed relevant publications on debugging strategies
- ❖ forward and backward snowballing from seven seed publications
- ❖ → 74 publications, 12 different debugging strategies

- ❖ different levels of abstraction
- ❖ contradictory definitions
- ❖ strategies that often apply only in specific study setups
- ❖ strategies that resemble other strategies, which makes them hard to distinguish

Which debugging strategies are helpful, which hinder
debugging performance bugs?

User Study

12 professional software developers
debugging a configuration-dependent performance bug

Density Converter

(multi platform image density conversion)

```
private static <T extends DensityDescriptor> Map<T, Dimension> getDensityBucketsWithFractionScale
    java.util.List<T> densities, Dimension srcDimension, Arguments args, float fraction) {
    double baseWidth = (double) srcDimension.width / fraction;
    double baseHeight = (double) srcDimension.height / fraction;
```

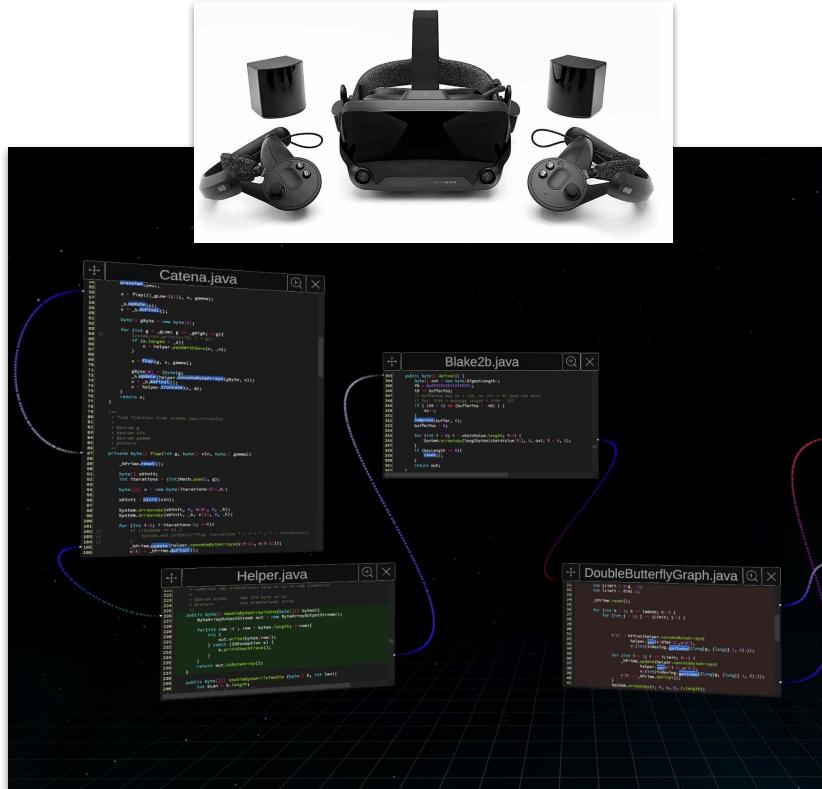
```
public final class Convert {
    public static void main(String[] rawArgs) {
        float fraction = 0.5f;
```



```
private void verticalFromWorkToDst(byte[][][] workPixels, byte[] outPixels, int start, int delta) {
    // ...
    for (int x = start; x < dstWidth; x += delta) {
        final int xLocation = x * nrChannels;
        for (int y = dstHeight - 1; y >= 0; y--) {
            final int max = verticalSubsamplingData.arrN[y];
            for (int j = max - 1; j >= 0; j--) {
                int valueLocation = verticalSubsamplingData.arrPixel[index];
                float arrWeight = verticalSubsamplingData.arrWeight[index];
                // ...
            }
        }
    }
}
```



SoftVR, Immersive Virtual Reality Debugging Tool



- ❖ position unlimited code windows with profiling information
- ❖ visualize control flow
- ❖ users can move through code-window setup

VR as novel study method:

- ❖ directly observe which code sections developers focus on
- ❖ analyze code-window patterns developers create

General Results

SoftVR is usable for performance bug detection (confirmed by SUS)

All developers reached the location of the bug in the code

7 developers successfully identified and reasoned about the bug

5 developers overlooked the bug or drew incorrect conclusions

Debugging as Episodes

Participants



Finding: Focus search and reasoning,
overcome episodes of aimlessness fast.

P11, reasoning:

"I need to understand what is going on in order to be able to formulate hypotheses".

P4, (problem-specific) exploration:

"Usually I try to get a fairly complete picture of something first."

P9, aimlessness:

[...] there is no point in staying in such a frustrating moment."

Reasoning appears substantially more often in successful debugging.

Aimlessness seems to be a possible factor for failure.

1

Assessing Energy Consumption



How can we **measure** software energy consumption?

Can we find a **reliable proxy**?

How can we get **code-level insights**?

Energy Consumption of Configurable Software Systems

2

Debugging Energy consumption



How can we identify **energy bugs**?



3

Practitioners' perspective



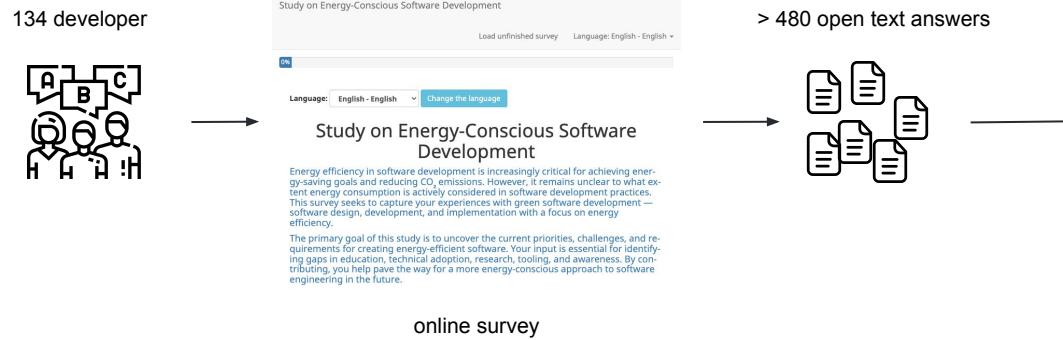
What is the **current state** of green software engineering?

Optimizing Software Energy Consumption in Practice



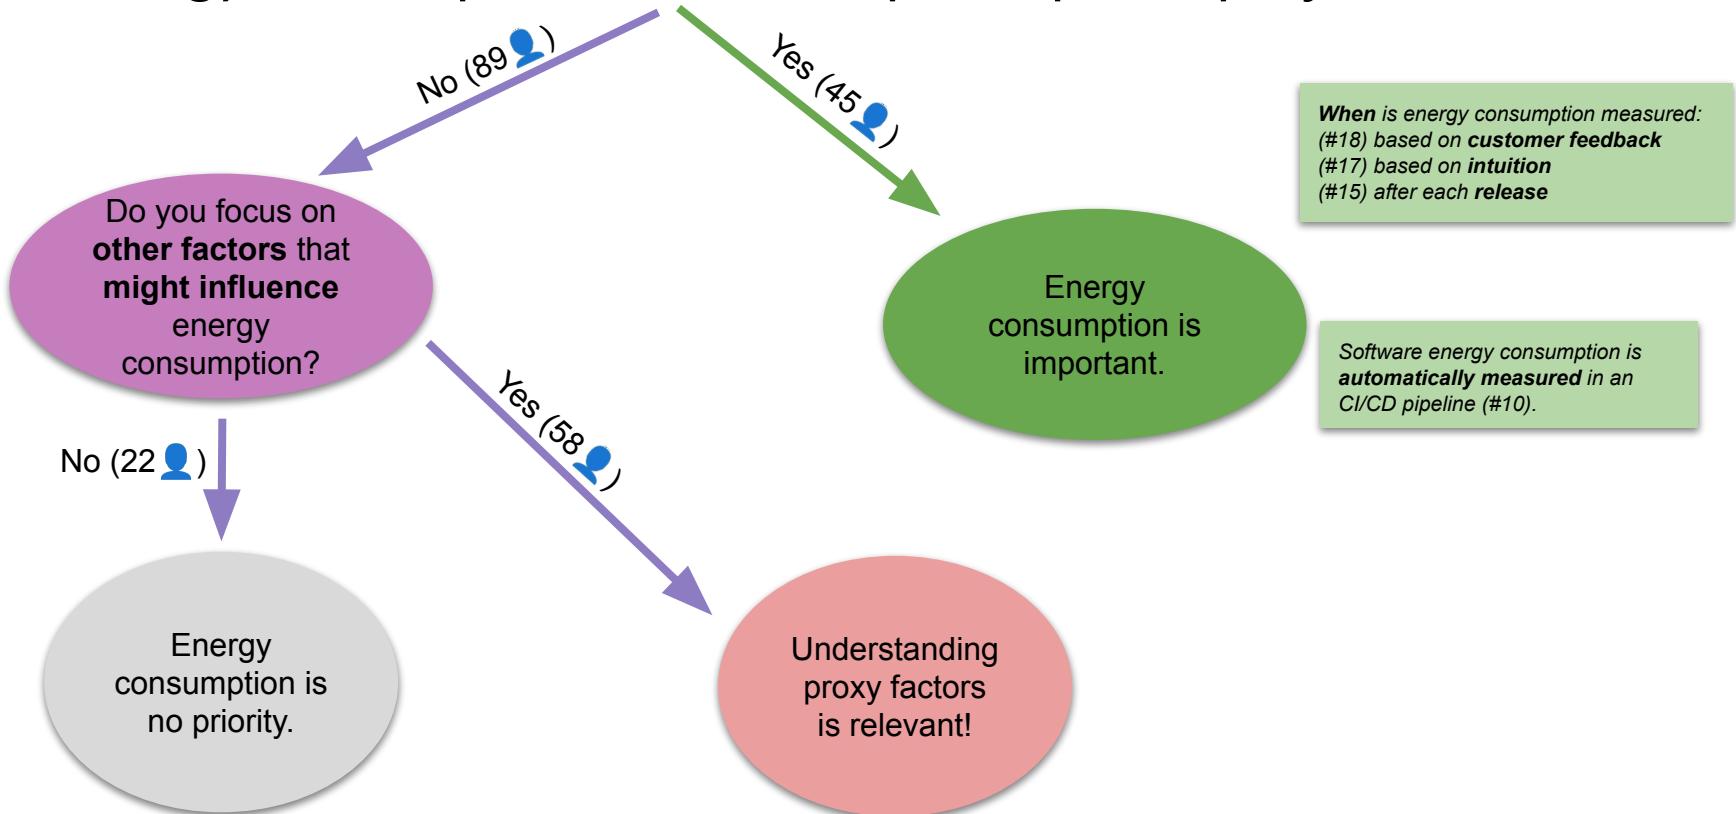
What is the **current state** of software energy consumption in practice?

How do practitioners **relate software energy consumption to other metrics** such as runtime performance?

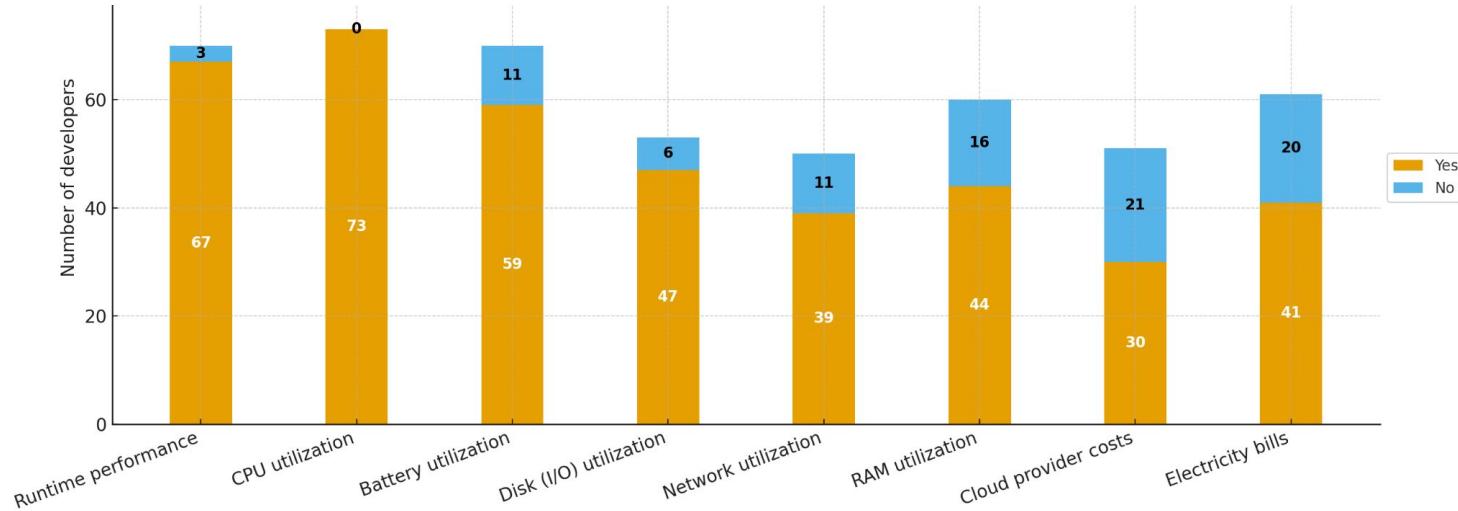


>16 hours manual analysis

Is energy consumption relevant in participants' projects?



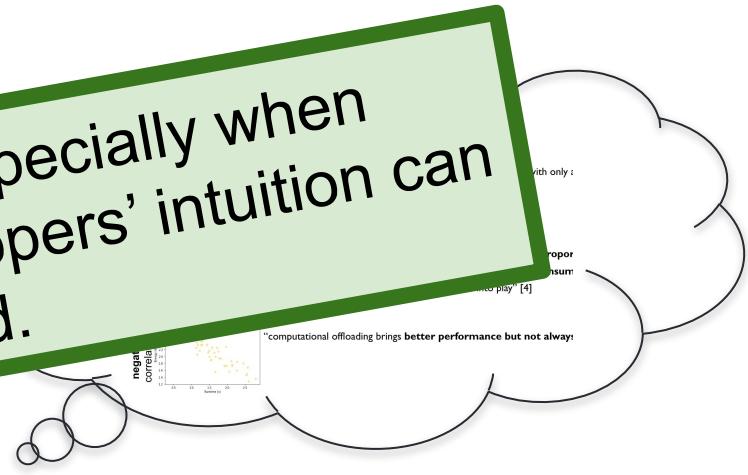
What metrics are considered good proxies for energy?



Do Developers Intuitively Know when Performance is a good Proxy for Energy?



Finding: It depends! Especially when positively correlated developers' intuition can be trusted.



How to Support Practitioners

Barriers to Overcome

Barriers preventing energy measurement:

- (55) Management is not interested
- (46) Customers are not interested
- (36) Not applicable in our use case

Persistent obstacles:

- (14) Lack of energy management knowledge
- (10) Energy efficiency is not a priority
- (06) Lack of education

Finding: It depends! Especially when positively correlated developers' intuition can be trusted.

Contra there are problems:

- (11) Energy consumption is not important
- (06) No perceived challenges
- (05) Energy is not considered a cost driver

Actions

(13) Increasing

er energy

Efficient
languages (C, C++, Rust)

(06) Introduce public regulations

(04) Enforce transparency,
especially for cloud providers

(05) Adopt efficient algorithms
and implementations



UNIVERSITÄT
LEIPZIG

Feel free to contact me!



max.weber@cs.uni-leipzig.de



<https://github.com/AI-4-SE>



linkedin.com/in/max-weber-84a83415b

