



UNIVERSITÄT
LEIPZIG

Green Configuration

Energy Consumption of Configurable Software Systems



Max Weber



Alina Mailach



Johannes Dorn



Christian Kalten.



Florian Sattler



Sven Apel



Norbert Siegmund

Energy Consumption at ecoCompute



Workloads, Use Cases

Software Systems

OS, Languages

Hardware



Rerouting the Cloud:
Cutting 90% CO₂ and
Cloud Costs with Smarter
Workload Shifting

Dryden Williams

CarbonRunner



Coding Smarter for a
Greener Future | A
comparison between Go
& Java

Moritz Böller

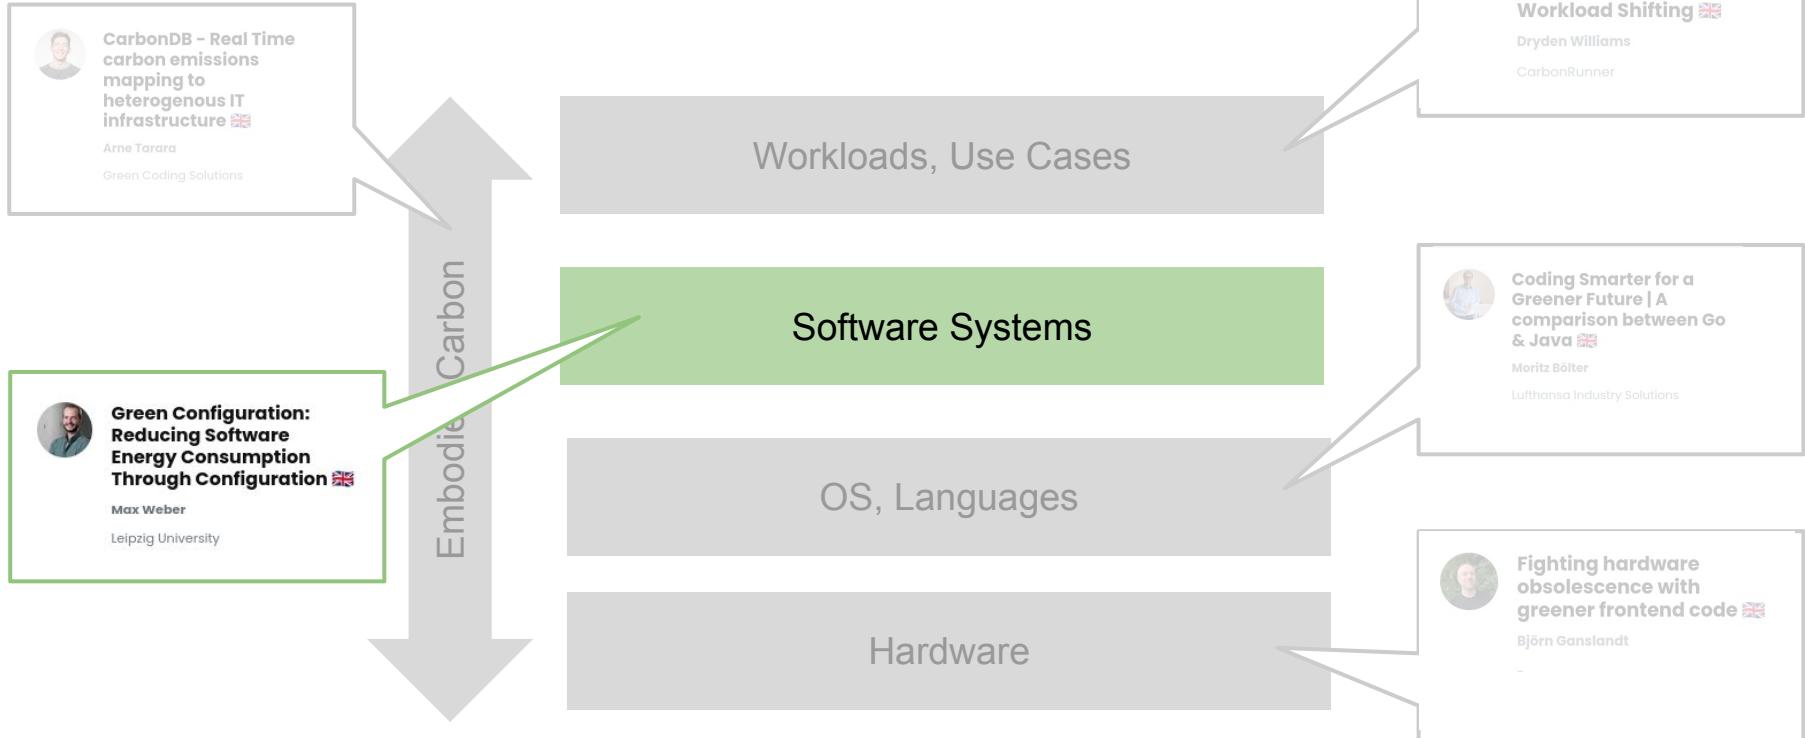
Lufthansa Industry Solutions



Fighting hardware
obsolescence with
greener frontend code

Björn Ganslandt

Energy Consumption at ecoCompute

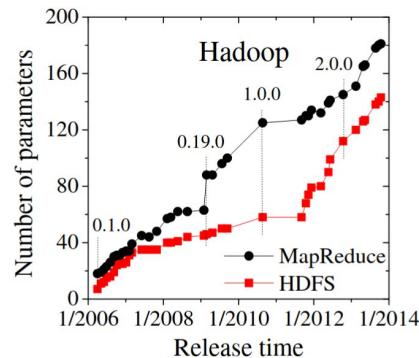
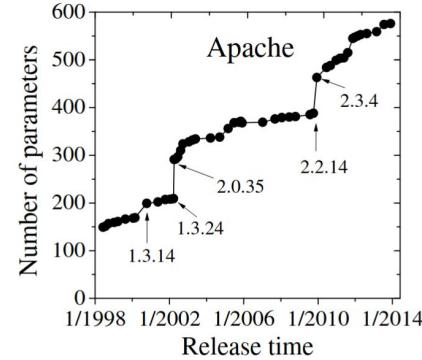
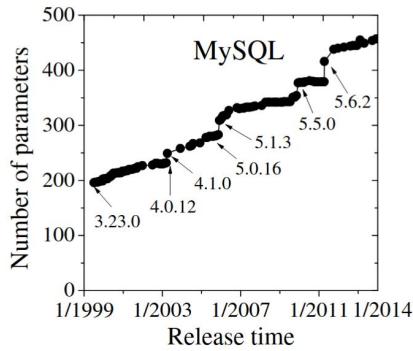


Energy Consumption at ecoCompute

How can we reduce the energy consumption of software systems?

By tailoring software to its workload and usage scenario.

Configuration is everywhere!



Number of parameters today:

MySQL – 700+

Apache – 900+

Hadoop – 400+

Users stick to default configurations,
leaving much potential unused.

1

Assessing Energy Consumption



How can we **measure** software energy consumption?

2

Debugging Energy consumption



3

Practitioners' perspective



Energy Consumption of Configurable Software Systems

Measuring Energy Consumption up to Code Level

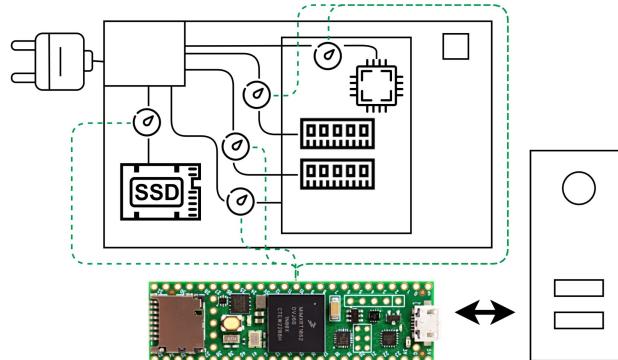


component-wise measurements
(adding up to the whole system)

High sampling rate:
3.5 kHz

no energy measurement bias

Accuracy: < 1.2%
measurement error



Hardware Components

Raspberry Pi 4B



60€

USB

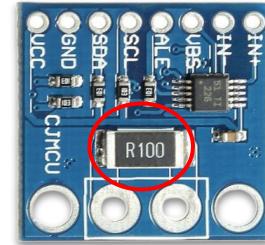
Teensy 4.1



30€

I²C

INA 226
Sample Time 280 us
Sample Rate 3.5 kHz



5€

3,500 measurements per second
≈ 1 sample every 0.3 ms (for all components)

175€

Hardware Components

Raspberry Pi 4B



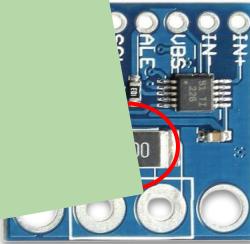
60€

USB

Teensy 4.1

Time spent until it worked reliably:
Priceless.

INA 226
Sample Time 280 us
Sample Rate 3.5 kHz

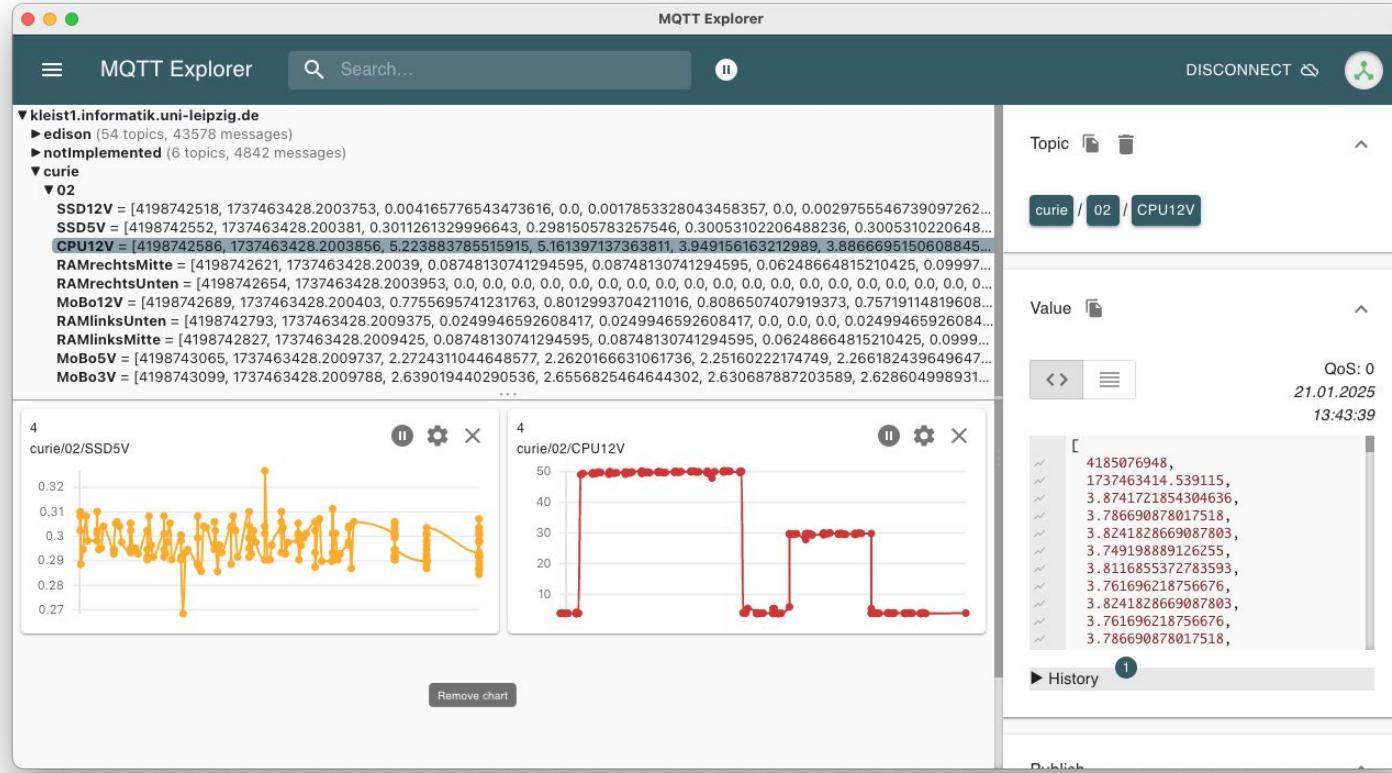


5€

3,500 measurements per second
≈ 1 sample every 0.3 ms (for all components)

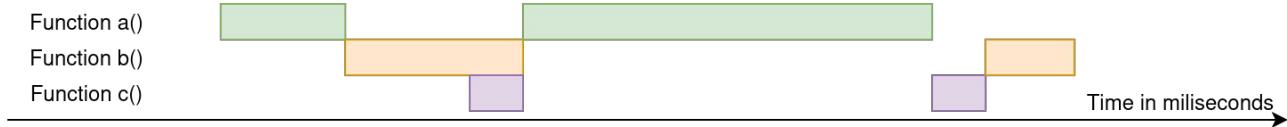
175€

Power Measurement

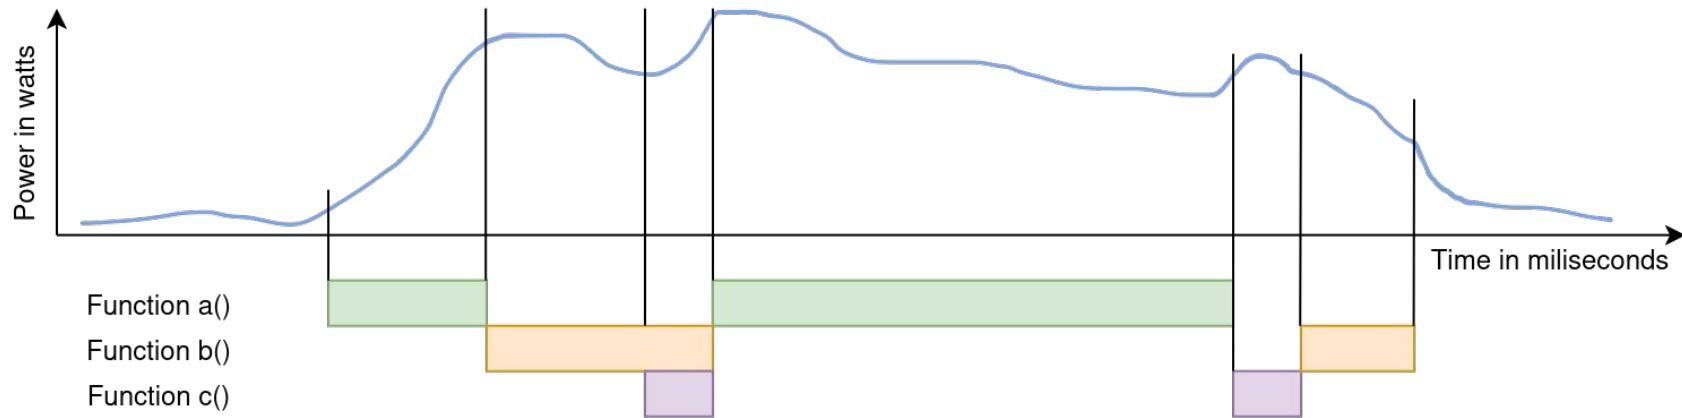


Performance Profiling

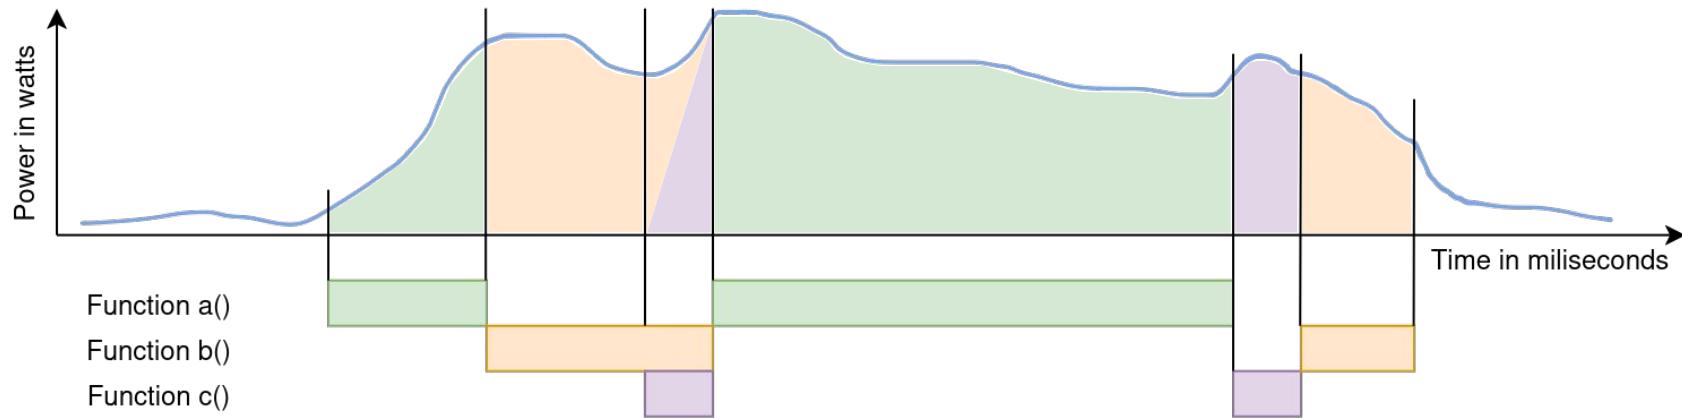
Linux *perf*



Synchronizing Energy and Performance Measurements



Synchronizing Energy and Performance Measurements



1

Assessing Energy Consumption



How can we **measure** software energy consumption?

Can we find a **reliable proxy**?

Energy Consumption of Configurable Software Systems

2

Debugging Energy consumption



3

Practitioners' perspective



Energy Metering is expensive, difficult and tool intensive.



Infeasible in many project settings.

Solution: Using performance as a proxy for energy consumption



System-Level Correlation

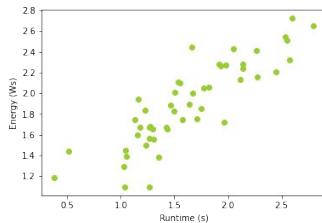


Option-Level Correlation



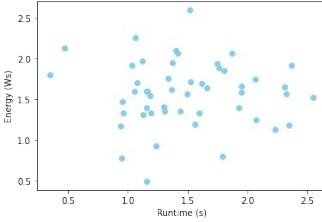
Literature Study Statements

positive correlation



“perform significantly better and consume less energy with only a small loss in QoS” [1]

no / absent correlation

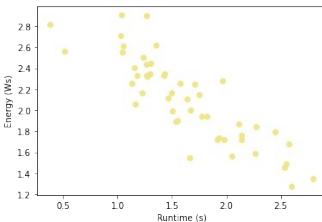


“the involved physical components may not consume energy proportional to time” [2]

“caching has an effect on performance but not on energy consumption” [2]

“no correlation if some communication comes into play” [4]

negative correlation



“computational offloading brings better performance but not always better energy efficiency” [3]

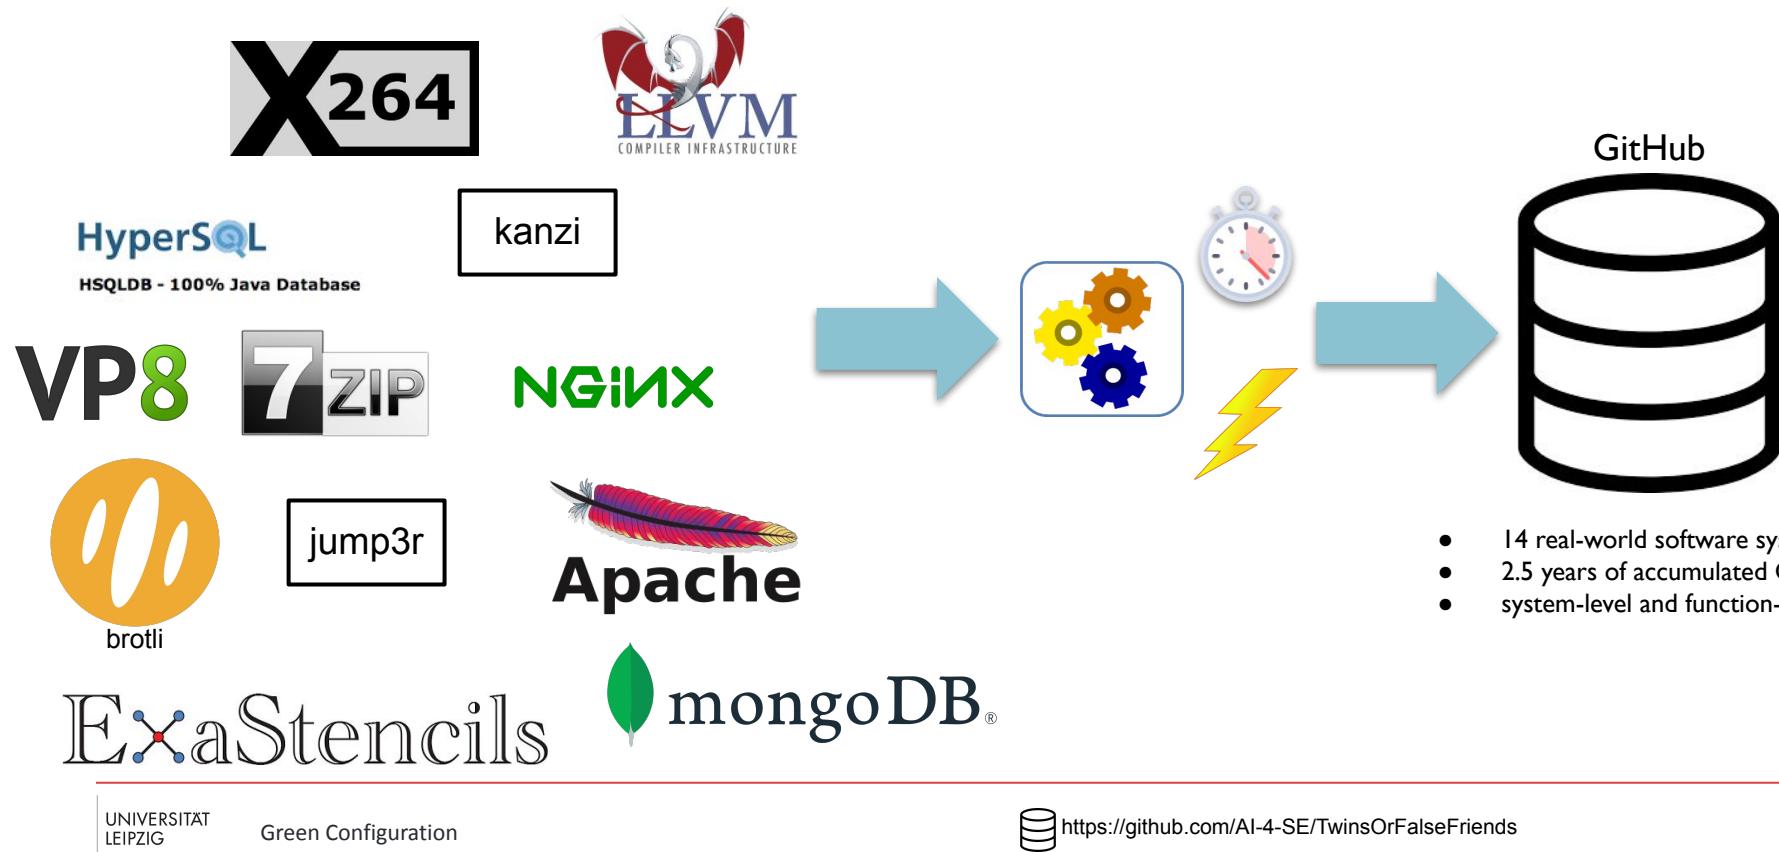
[1] Green: a framework for supporting energy-conscious programming using controlled approximation, 2020

[2] Evaluating the Impact of Caching on the Energy Consumption and Performance of Progressive Web Apps, 2020

[3] Analysis of Performance and Energy Consumption of Wearable Devices and Mobile Gateways in IoT Applications, 2019

[4] Quantifying energy use in dense shared memory HPC node, 2016

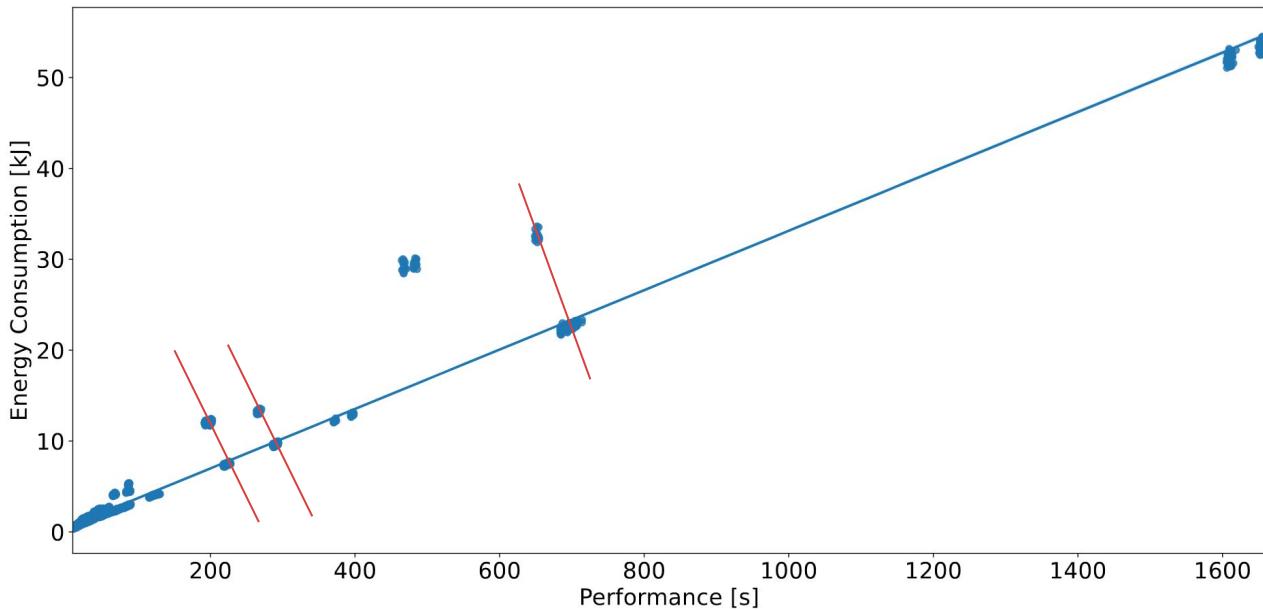
Empirical Study to Investigate Performance-Energy Correlation



System-Level Correlation

lzip

Correlation: 0.99



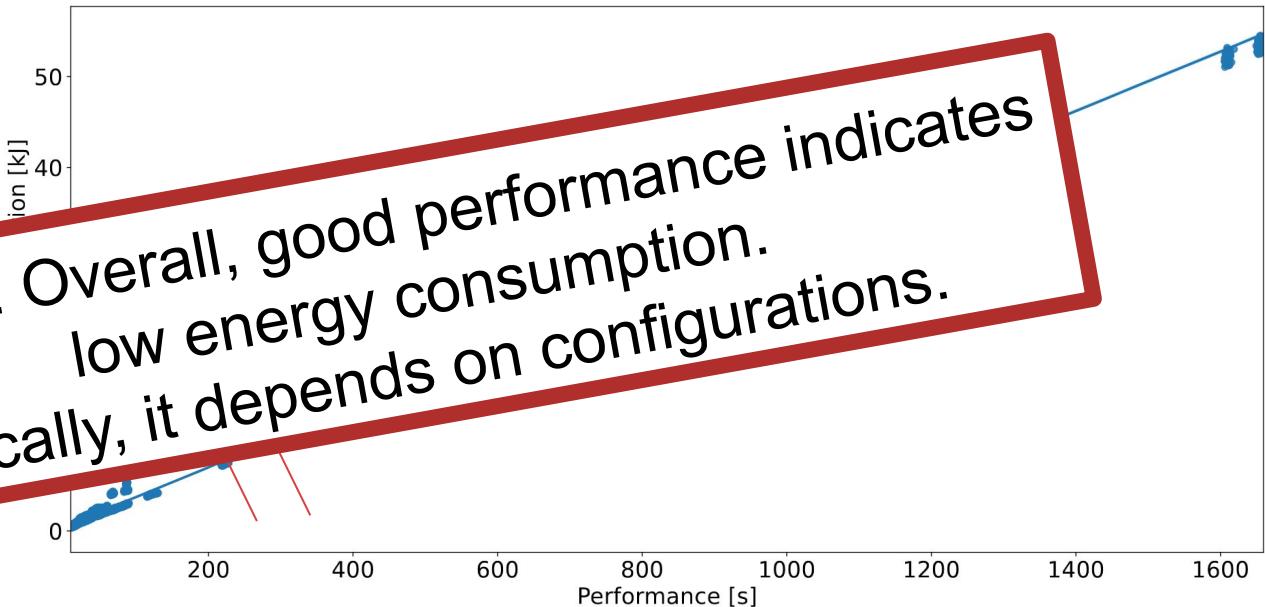
System-Level Correlation



Correl.

Irzip

Finding: Overall, good performance indicates low energy consumption.
Locally, it depends on configurations.



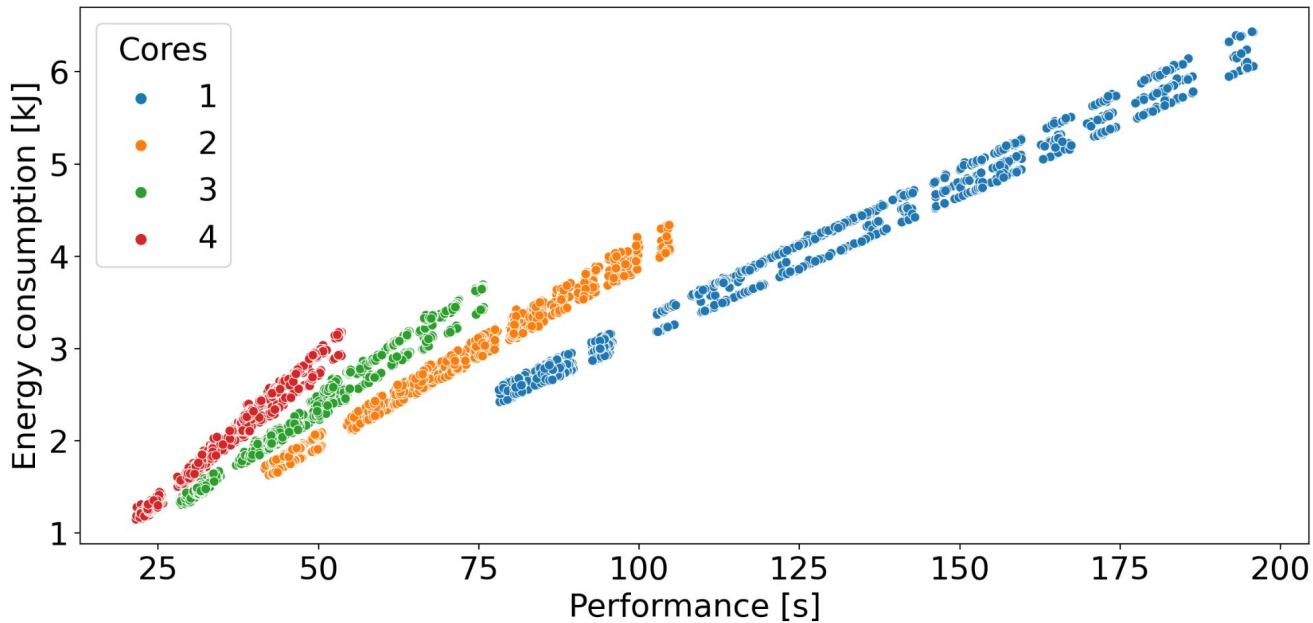
Option-Level Correlation



Correlation: 0.99
MAPE: 10.3%



Correlation: 0.99
MAPE: 2.5%



Option-Level Correlation

HyperSQL

HSQLDB - 100% Java Database



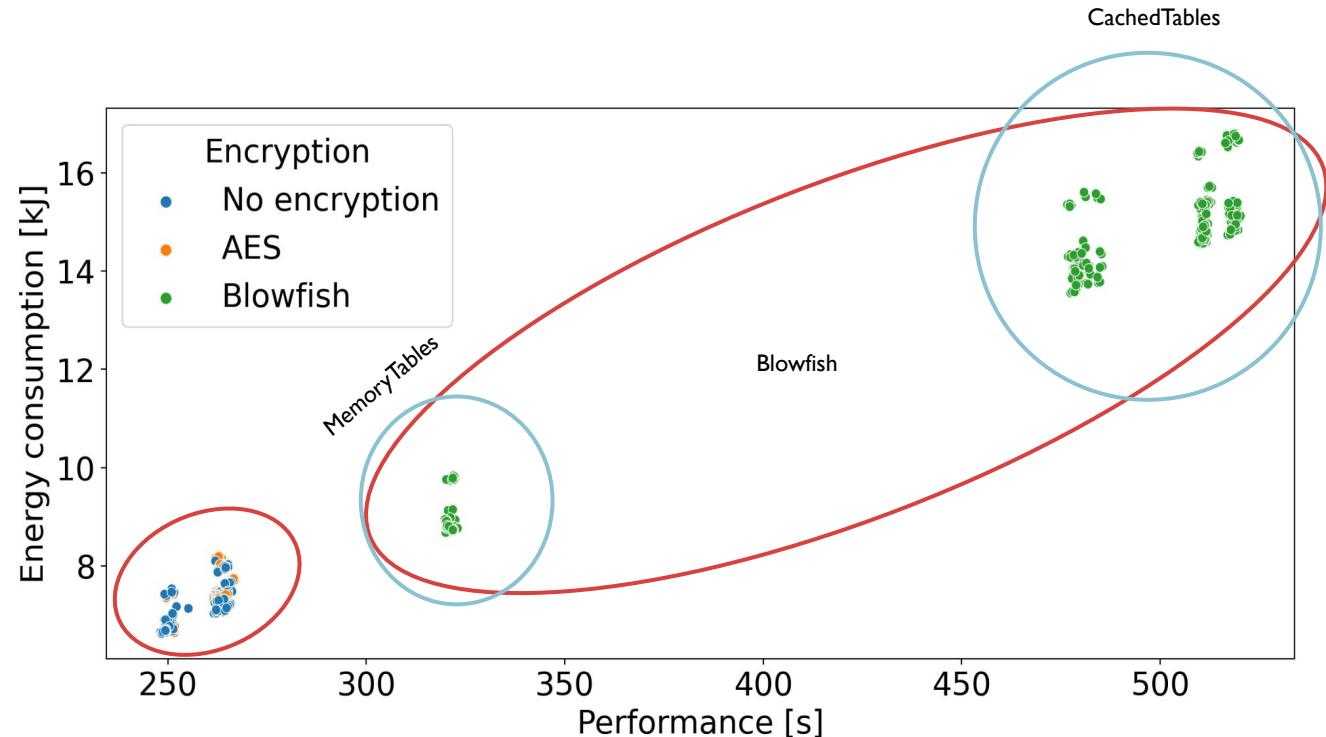
Correlation: 0.99
MAPE: 3.0%



Blowfish & MemoryTables



Correlation: 0.13
MAPE: 213.5% !



Option-Level Correlation

HyperSQL

HSQLDB - 100% Java Database



Correlation:
MAPE:

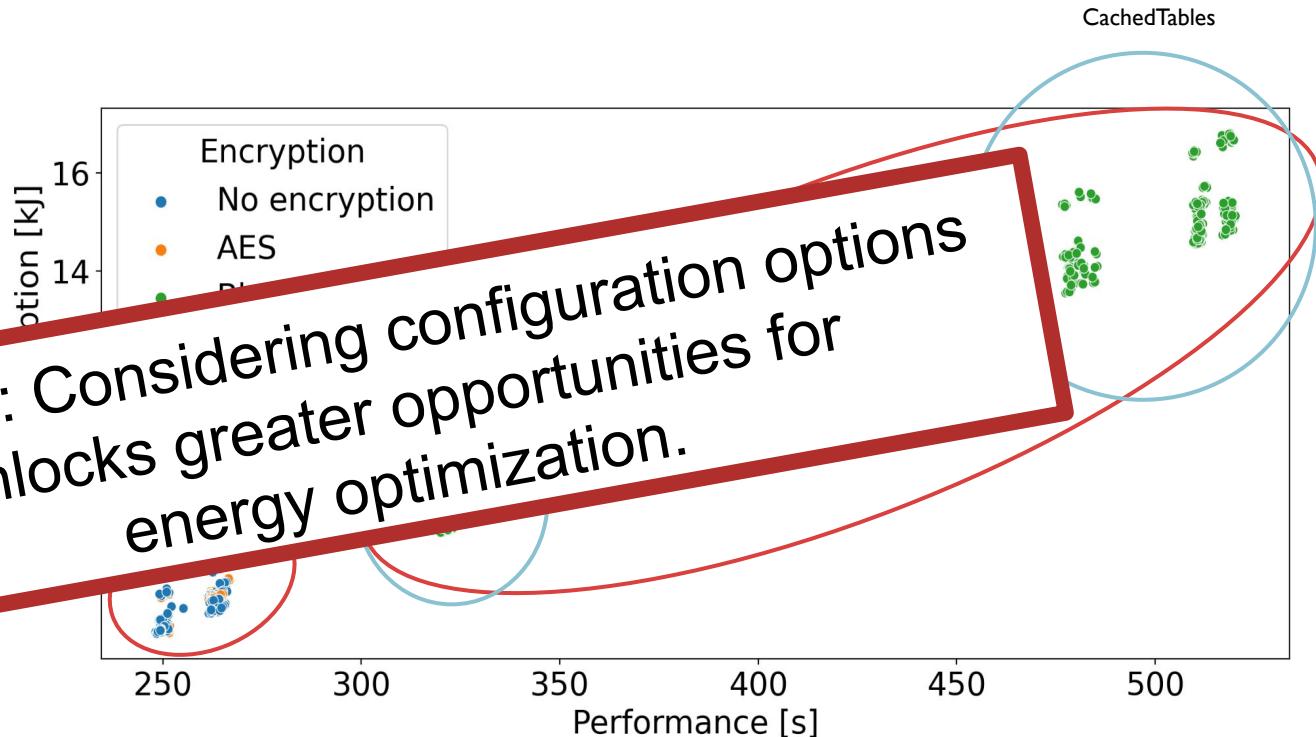


Finding: Considering configuration options
unlocks greater opportunities for
energy optimization.

Blowfish & Memory Tables



Correlation: 0.13
MAPE: 213.5%



1

Assessing Energy Consumption



How can we **measure** software energy consumption?

Can we find a **reliable proxy**?

How can we get **code-level insights**?

Energy Consumption of Configurable Software Systems

2

Debugging Energy consumption

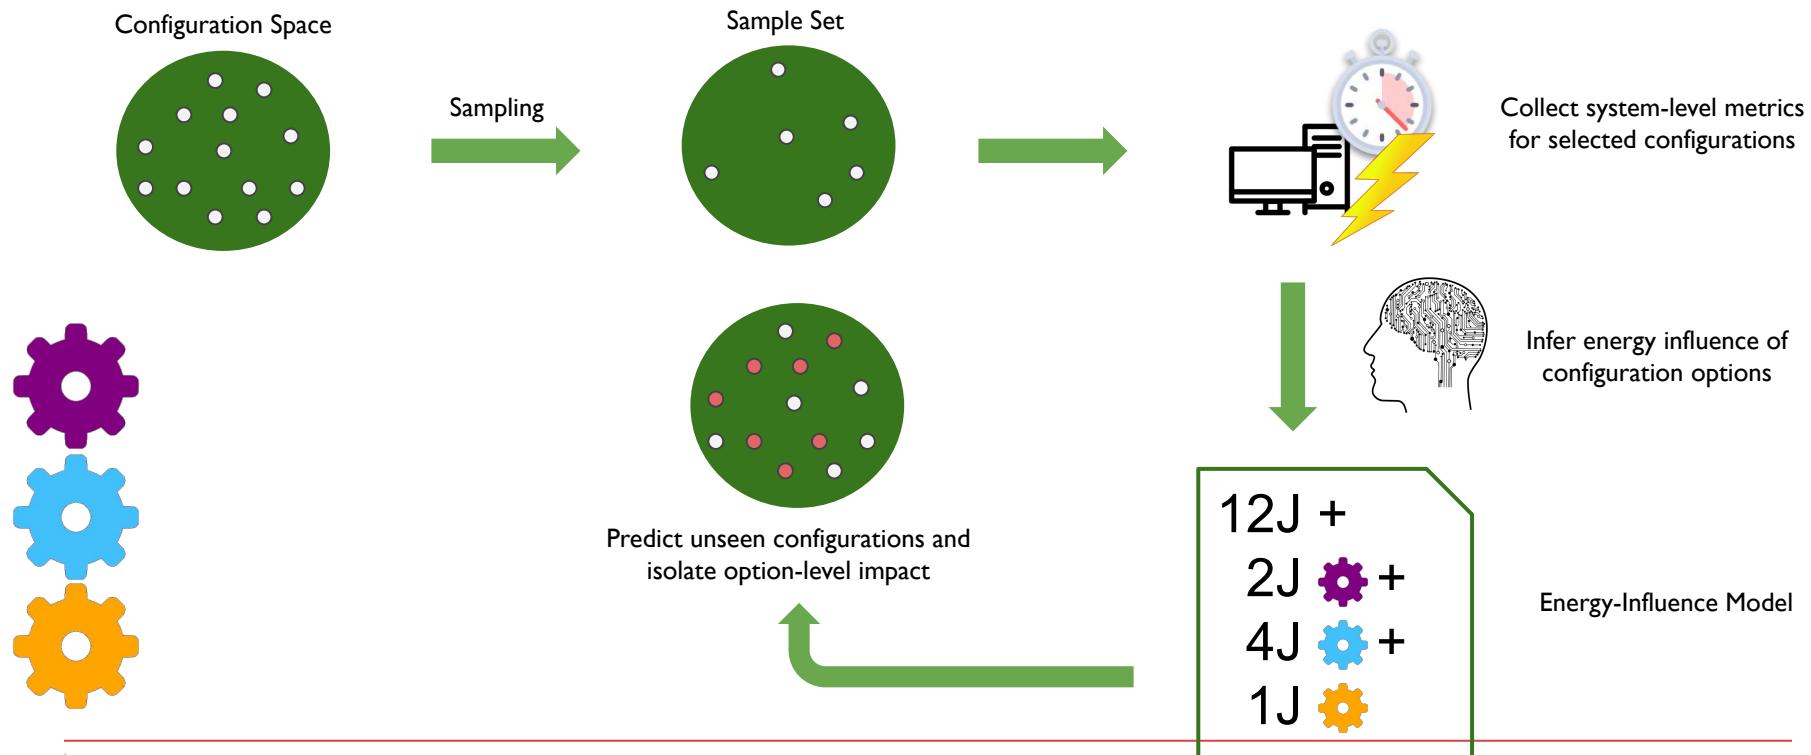


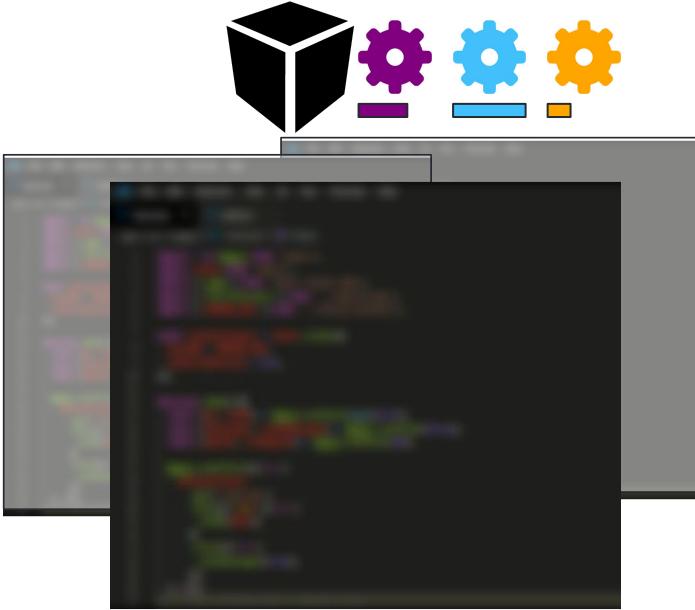
3

Practitioners' perspective



Building Energy-Influence Models for Configurable Systems



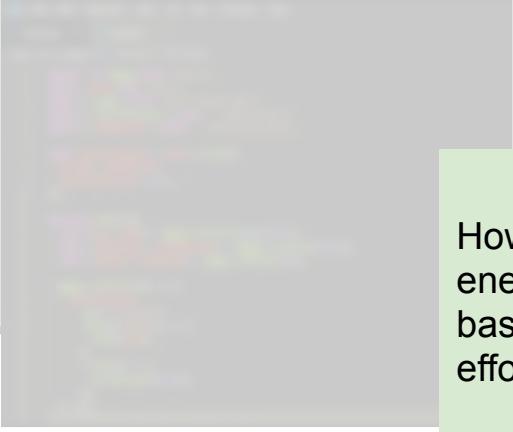
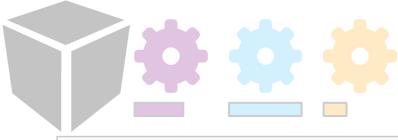


As a **developer**, how do I know which part of my code I need to change?

Black-Box Models

As a **user**, I can optimize configurations to reduce energy consumption 💪

... but I can't see **where** in the code the energy goes.

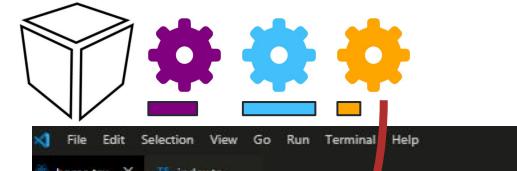


How can we **efficiently** pin point energy hotspots in the code base with minimal measurement effort?

Black-Box Models

As a **user**, I can optimize configurations to reduce energy consumption 

... but I can't see **where** in the code the energy goes.



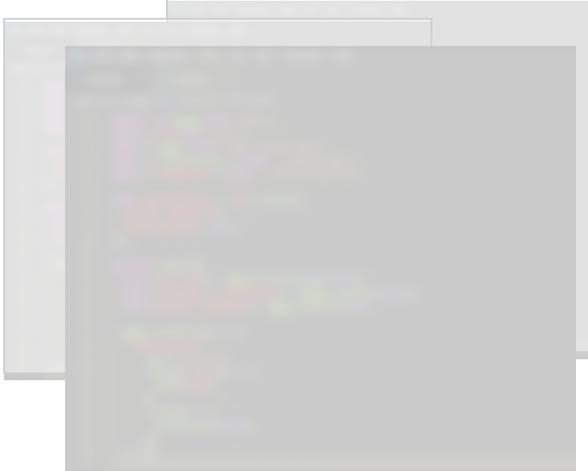
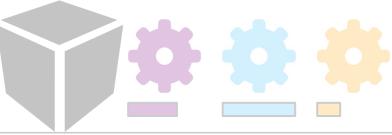
A screenshot of a code editor showing a file named `home.tsx`. The code is a React component that uses axios to fetch user data from a server. A red arrow points from the top diagram to the line of code where `axiosInstance.get('/user/me')` is called.

```
1 import * as React from 'react';
2 import axios from 'axios';
3 import { Link } from 'react-router-dom';
4 import { startCheckout } from './lib/stripe';
5 import { SERVER_URL } from '../utils/constants';
6
7 const axiosInstance = axios.create({
8   baseURL: SERVER_URL,
9   withCredentials: true,
10 });
11
12 function Home() {
13   const [me, setMe] = React.useState<any>(null);
14   const [showLogin, setShowLogin] = React.useState(false);
15   const [amount, setAmount] = React.useState(10);
16
17   React.useEffect(() => {
18     axiosInstance
19       .get('/user/me')
20       .then((data) => {
21         setMe(data);
22       })
23       .catch(() => {
24         setShowLogin(true);
25       });
26   }, []);
}
```

Established White-Box Models

We know exactly **which** code statements consume energy.

...but measuring this across a full system is often **infeasible** 😞



Black-Box Models

As a **user**, I can optimize configurations to reduce energy consumption 💪

... but I can't see **where** in the code the energy goes. 😞

```

1 import * as React from 'react';
2 import axios from 'axios';
3 import { Link } from 'react-router-dom';
4 import { startCheckout } from '../lib/stripe';
5 import { SERVER_URL } from './utils/constants';
6
7 const axiosInstance = axios.create({
8   baseURL: SERVER_URL,
9   withCredentials: true,
10 });
11
12 function Home() {
13   const [me, setMe] = React.useState<any>(null);
14   const [showLogin, setShowLogin] = React.useState(false);
15   const [amount, setAmount] = React.useState(10);
16
17   React.useEffect(() => {
18     axiosInstance
19       .get('/user/me')
20       .then(({ data }) => {
21         setMe(data);
22       })
23       .catch(() => {
24         setShowLogin(true);
25       });
26   }, []);

```

Method-Level White-Box Models

As a user, I still benefit from efficient configurations that save energy ⚡.

As a developer, I can identify which methods consume energy without measuring every single line of code.

```

1 import * as React from 'react';
2 import axios from 'axios';
3 import { Link } from 'react-router-dom';
4 import { startCheckout } from '../lib/stripe';
5 import { SERVER_URL } from './utils/constants';
6
7 const axiosInstance = axios.create({
8   baseURL: SERVER_URL,
9   withCredentials: true,
10 });
11
12 function Home() {
13   const [me, setMe] = React.useState<any>(null);
14   const [showLogin, setShowLogin] = React.useState(false);
15   const [amount, setAmount] = React.useState(10);
16
17   React.useEffect(() => {
18     axiosInstance
19       .get('/user/me')
20       .then(({ data }) => {
21         setMe(data);
22       })
23       .catch(() => {
24         setShowLogin(true);
25       });
26   }, []);

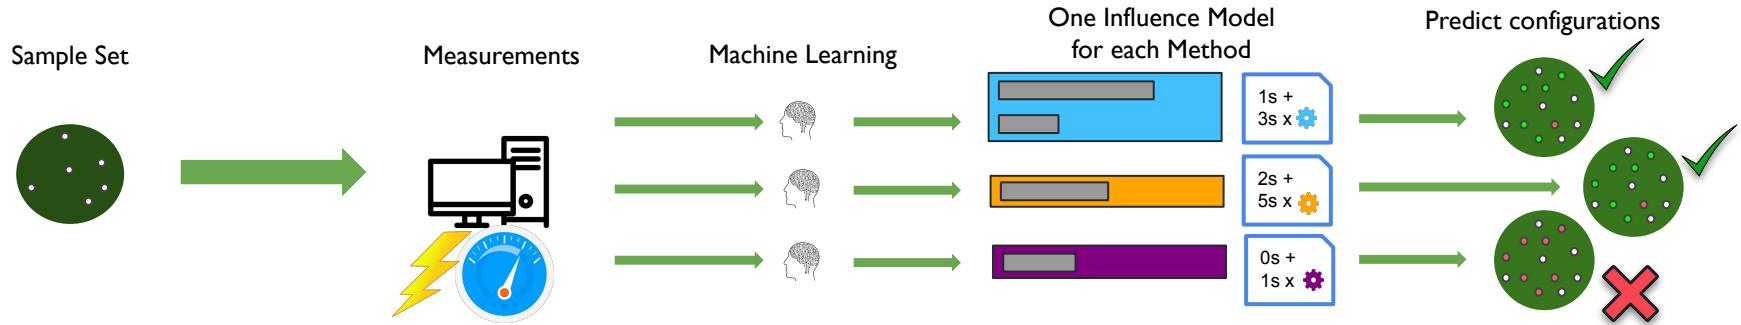
```

Established White-Box Models

We know exactly **which** code statements consume energy. 💪

...but measuring this across a full system is often **infeasible**. 😞

White-Box Energy Influence Models via Profiling



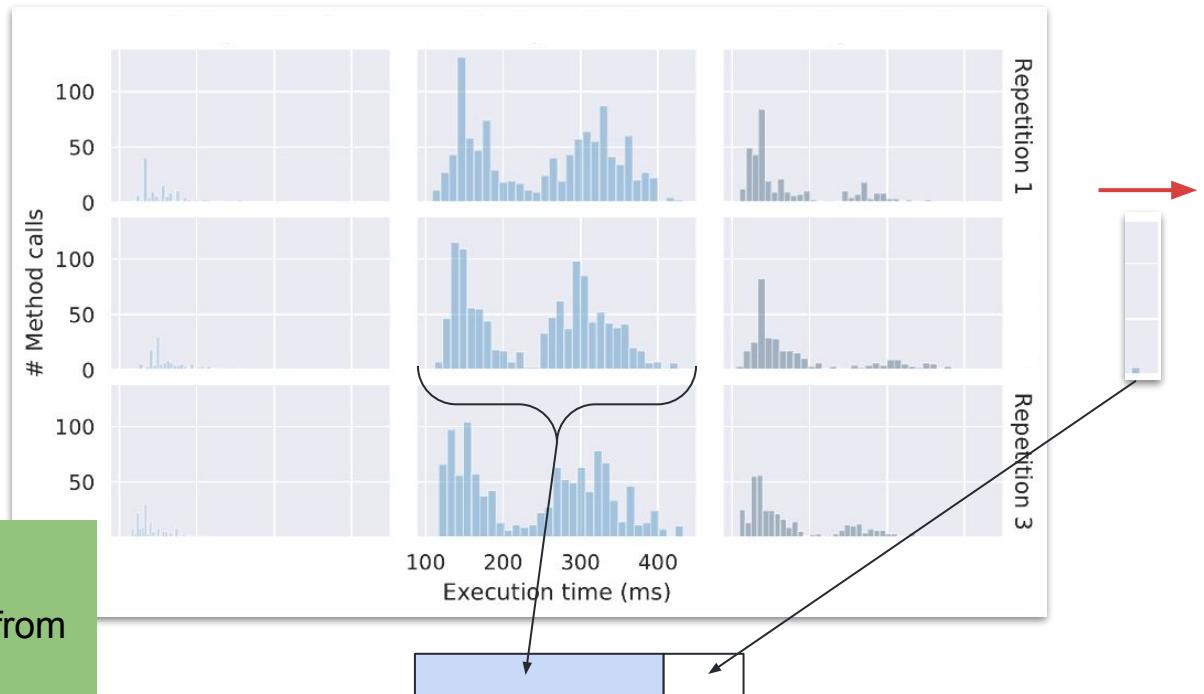
Challenge: Handling Variance in Measurement Data

Context variance

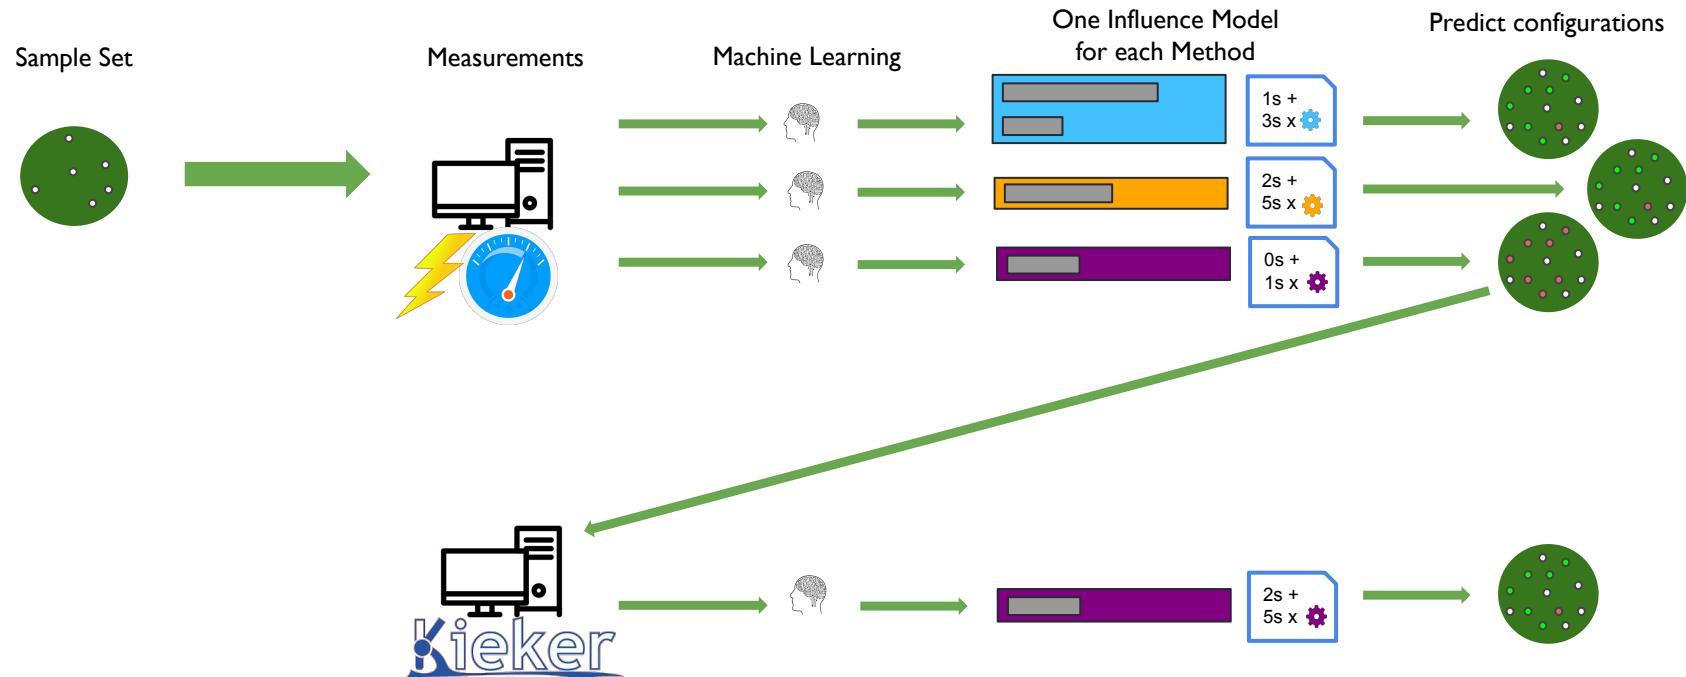
Configuration variance

Measurement variance

Solution: Filter large,
non-deterministic outliers from
context variance.



White-Box Energy Influence Models via Profiling



IDE Integration

The screenshot shows the Eclipse Platform interface with the title "runtime-EclipseApplication - catena/src/main/java/components/graph/algorithms/GenericGraph.java - Eclipse Platform".

Package Explorer: Shows the project structure under "catena [GreenDev master]". The file "DoubleButterflyGraph.java" is selected.

Code Editor: Displays the Java code for "GenericGraph.java". The code implements a graph algorithm using byte arrays and helper objects.

```
public byte[][] graph (int g, byte[][] v, int lambda){  
    int dim1 = (int) Math.pow(2, g);  
    int dim2 = hPrime.getOutputSize();  
    byte[][] r = new byte[dim1][dim2];  
  
    for (int k = 0; k < lambda; ++k){  
        r[0] = hFirst(helper.concatenateByteArrays(v[dim1-1],  
            v[indexing.getIndex(0, g)]));  
  
        int loop = (int) Math.pow(2, g);  
  
        for (int i = 1; i < loop; ++i){  
            hPrime.update(helper.concatenateByteArrays(r[i-1], v[indexing.getIndex(i, g)]));  
            r[i] = hPrime.doFinal();  
        }  
  
        System.arraycopy(r, 0, v, 0, r.length);  
    }  
    return v;  
}  
  
private byte[] hFirst(byte[] in){  
    int n = h_.getOutputSize();  
    int k = hPrime.getOutputSize();  
    int l = k/n;  
  
    byte[][] w = new byte[l][n];  
    byte[] iByte = new byte[l];  
}
```

Performance Overview: A table showing performance values and paths for various components.

Performance Value	Path
100.0	main.java.components.graph.algorithms.GenericGraph:graph
1.061113	main.java.Catena:flap
0.19015925	profiling.CatenaProfiler:initCatenaByConfig
0.1401815	main.java.components.gamma.algorithms.SaltMix:gamma

1

Assessing Energy Consumption



How can we **measure** software energy consumption?

Can we find a **reliable proxy**?

How can we get **code-level insights**?

Energy Consumption of Configurable Software Systems

2

Debugging Energy consumption



How do developers identify **energy bugs**?



3

Practitioners' perspective



Energy Bugs in Configurable Systems

Functional bug



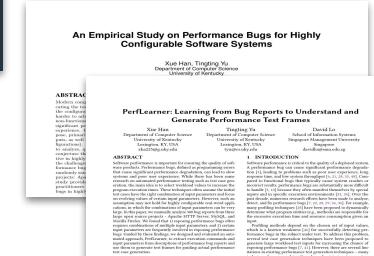
Energy bugs



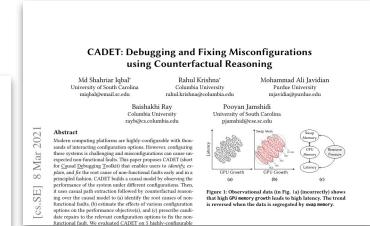
Configuration-dependent bugs



Requires fixing code



Fixes are time consuming



What Does Debugging Literature Tell Us?

Software debugging, testing, and verification

by B. Helpern
P. Santhanam

In commercial software development organizations, increased complexity of products, shortening of development cycles, and customer expectations of quality have placed a major responsibility on the developer for debugging, testing, and verification. As this issue of the IBM Systems Journal illustrates, there are numerous approaches to the problem.

in cost to develop a programming system product from an isolated program (see Figure 1).

With the advent of the Internet and the World Wide Web, the problems that were recognized a quarter century ago as having "no silver bullet" for a solution are now more pressing than ever.

A systematic literature review on benchmarks for evaluating debugging approaches^{*}

Thomas Hirsch, Birgit Hofer^{*}
Institute of Software Technology, Graz University of Technology, Austria

ARTICLE INFO

Article history:
Received 16 December 2021
Received in revised form 22 June 2022
Accepted 23 June 2022
Available online 30 June 2022

Keywords:
Debugging
Benchmark

ABSTRACT

Bug benchmarks are used in development and evaluation of debugging and automation regard. Quantitative performance comparison is only possible when they have been evaluated on the same benchmarks are often specialized towards usage for certain debug data, metrics, and artifacts. Such benchmarks cannot be easily used as scope as such approach may rely on specific data such as are not included in the dataset. Furthermore, benchmarks vary in subject programs and the size of the individual subject pro-

- ❖ reviewing 74 publications
- ❖ revealed 12 debugging strategies

On Debugging the Performance of Configurable Software Systems: Developer Needs and Tailored Tool Support

Miguel Velez
Carnegie Mellon University
Pooyan Jamshidi
University of South Carolina

Norbert Siegmund
Leipzig University
Sven Apel
Saarland Informatics Campus - Saarland University
Christian Kästner
Carnegie Mellon University

ABSTRACT
Configurable software is considered to be a software system that is a generic bug if it was misclassified as other challenges. While there are numerous debugging techniques that can support dev-



!

No clear consensus -
literature shows fragmented
perspectives on debugging.

[Miguel Velez, Pooyan Jamshidi, Norbert Siegmund, Sven Apel, and Christian Kästner. On debugging the performance of configurable software systems: Developer needs and tailored tool support]

Let's observe developers debugging energy bugs in VR!

A 3D wireframe cube grid is visible in the background, suggesting a virtual environment.

Five code editor windows are displayed on the wireframe cube:

- Catena.java**:

```
private void x_flop(ByteArrayInputStream x, byte[] gamma) {
    x.flip();
    x.read();
    x.update();
    x.unread();
}

byte[] byteToNewByteArray() {
    for (int i = 0; i < 100; i++) {
        if (x.length() == 0) {
            x = x.readUnderline(x, 0);
        }
        x.read();
    }
    return x;
}

// Flop function from catena specification
private void x_flop(ByteArrayInputStream x, byte[] gamma) {
    byte[] bytes = new byte[1];
    int iterations = (int) Math.pow(2, 0);
    bytes[0] = new byte[iterations][0];
    x.read(bytes[0]);
    System.arraycopy(bytes[0], 0, bytes, 0, 1);
    for (int i = 0; i < iterations; i++) {
        if (bytes[0][i] == 0) {
            bytes[0][i] = (byte) (x.read() / iterations);
        }
    }
    x.write(bytes);
}
```
- Blake2b.java**:

```
public byte[] digest() {
    byte[] out = new byte[digestLength];
    try {
        buffer.position(0);
        buffer.limit(128);
        if (digestLength == 128) {
            System.arraycopy(buffer, 0, out, 0, 128);
        } else {
            if ((out.length - 1) * 128 >= buffer.position()) {
                out[out.length - 1] = 0;
            }
        }
        digestBB(buffer, 0);
        System.arraycopy(buffer, 0, out, 0, digestLength);
    } catch (IOException e) {
        e.printStackTrace();
    }
    return out;
}
```
- Blake2b_1.java**:

```
public byte[] digest() {
    byte[] out = new byte[digestLength];
    try {
        buffer.position(0);
        buffer.limit(128);
        if (digestLength == 128) {
            System.arraycopy(buffer, 0, out, 0, 128);
        } else {
            if ((out.length - 1) * 128 >= buffer.position()) {
                out[out.length - 1] = 0;
            }
        }
        digestBB(buffer, 0);
        System.arraycopy(buffer, 0, out, 0, digestLength);
    } catch (IOException e) {
        e.printStackTrace();
    }
    return out;
}
```
- DoubleButterflyGraph.java**:

```
private void compress(ByteArrayOutputStream out, byte[] message, int messagePos) {
    int stackSize = 0;
    int items = 0;
    byte[] item = new byte[1];
    _vertex.reset();
    for (int k = 0; k < 1000; k++) {
        for (int i = 0; i < 1000; i++) {
            for (int j = 0; j < 1000; j++) {
                if (_vertex.get(i) & _vertex.get(j)) {
                    item[0] = 1;
                } else {
                    item[0] = 0;
                }
                out.write(item);
            }
        }
    }
}
```
- Helper.java**:

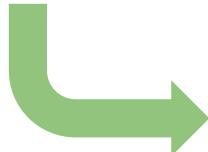
```
* CURRENT TWO DIMENSIONAL SYSTEM WORKS ON ONE DIMENSIONAL
444 * CURRENT TWO DIMENSIONAL SYSTEM WORKS ON ONE DIMENSIONAL
445 *
446 * update bytes
447 * one dimensional array
448 * public byte[] doubleByteArrayToInt(ByteArrayInputStream bytes) {
449 *     ByteArrayOutputStream out = new ByteArrayOutputStream();
450 *     for (int row = 0; row < bytes.length(); row++) {
451 *         for (int col = 0; col < bytes.length(); col++) {
452 *             autoUpdateBytes(out, bytes, row, col);
453 *             if (row == col) {
454 *                 autoUpdateBytes(out, bytes, col, col);
455 *             }
456 *         }
457 *     }
458 *     return out.toByteArray();
459 * }
460 * public byte[] oneDimByteArrayToInt(ByteArrayInputStream bytes) {
461 *     byte[] item = new byte[1];
462 *     for (int i = 0; i < bytes.length(); i++) {
463 *         item[0] = bytes.read();
464 *         out.write(item);
465 *     }
466 * }
```

Case Study: Density Converter

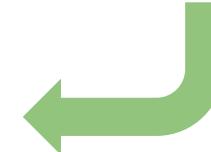
(multi platform image density conversion)

```
private static <T extends DensityDescriptor> Map<T, Dimension> getDensityBucketsWithFractionScale
    java.util.List<T> densities, Dimension srcDimension, Arguments args, float fraction) {
    double baseWidth = (double) srcDimension.width / fraction;
    double baseHeight = (double) srcDimension.height / fraction;
```

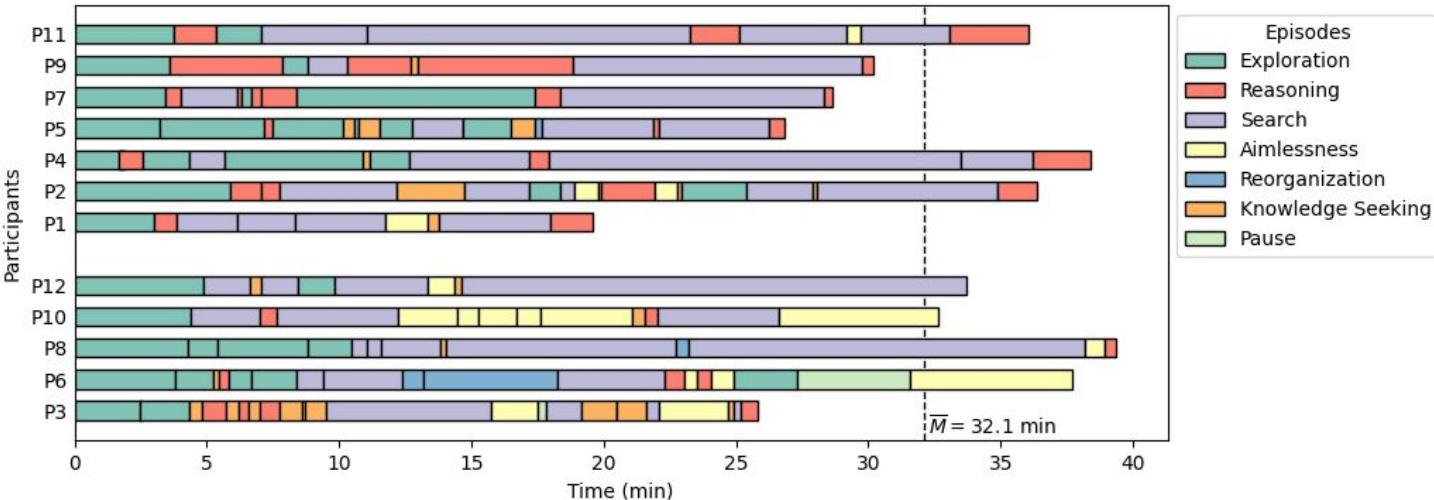
```
public final class Convert {
    public static void main(String[] rawArgs) {
        float fraction = 0.5f;
```



```
private void verticalFromWorkToDst(byte[][][] workPixels, byte[] outPixels, int start, int delta) {
    // ...
    for (int x = start; x < dstWidth; x += delta) {
        final int xLocation = x * nrChannels;
        for (int y = dstHeight - 1; y >= 0; y--) {
            final int max = verticalSubsamplingData.arrN[y];
            for (int j = max - 1; j >= 0; j--) {
                int valueLocation = verticalSubsamplingData.arrPixel[index];
                float arrWeight = verticalSubsamplingData.arrWeight[index];
            }
        }
    }
}
```



Debugging as Episodes



P11, reasoning:

"I need to understand what is going on in order to be able to formulate hypotheses".

P4, (problem-specific) exploration:

"Usually I try to get a fairly complete picture of something first."

P9, aimlessness:

"[...] there is no point in staying in such a frustrating moment."

Reasoning appears substantially more often in successful debugging.

Aimlessness seem to be a possible predictor for failure.

Debugging as Episodes



P11, reasoning:

"I need to understand what is going on in order to be able to formulate hypotheses".

P4, (problem-specific) exploration:

"Usually I try to get a fairly complete picture of something first."

P9, aimlessness:

[...] there is no point in staying in such a frustrating moment."

Reasoning appears substantially more often in successful debugging.

Aimlessness seems to be a possible factor for failure.

1

Assessing Energy Consumption



How can we **measure** software energy consumption?

Can we find a **reliable proxy**?

How can we get **code-level insights**?

2

Debugging Energy consumption



How can we identify **energy bugs**?



3

Practitioners' perspective



What is the **current state** of green software engineering?

Energy Consumption of Configurable Software Systems

Optimizing Software Energy Consumption in Practice



What is the **current state** of software energy consumption in practice?

How do practitioners **relate software energy consumption to other metrics** such as runtime performance?



Study on Energy-Conscious Software Development

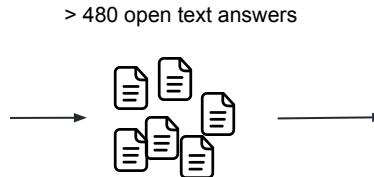
Load unfinished survey Language: English - English ▾

Language: English - English Change the language

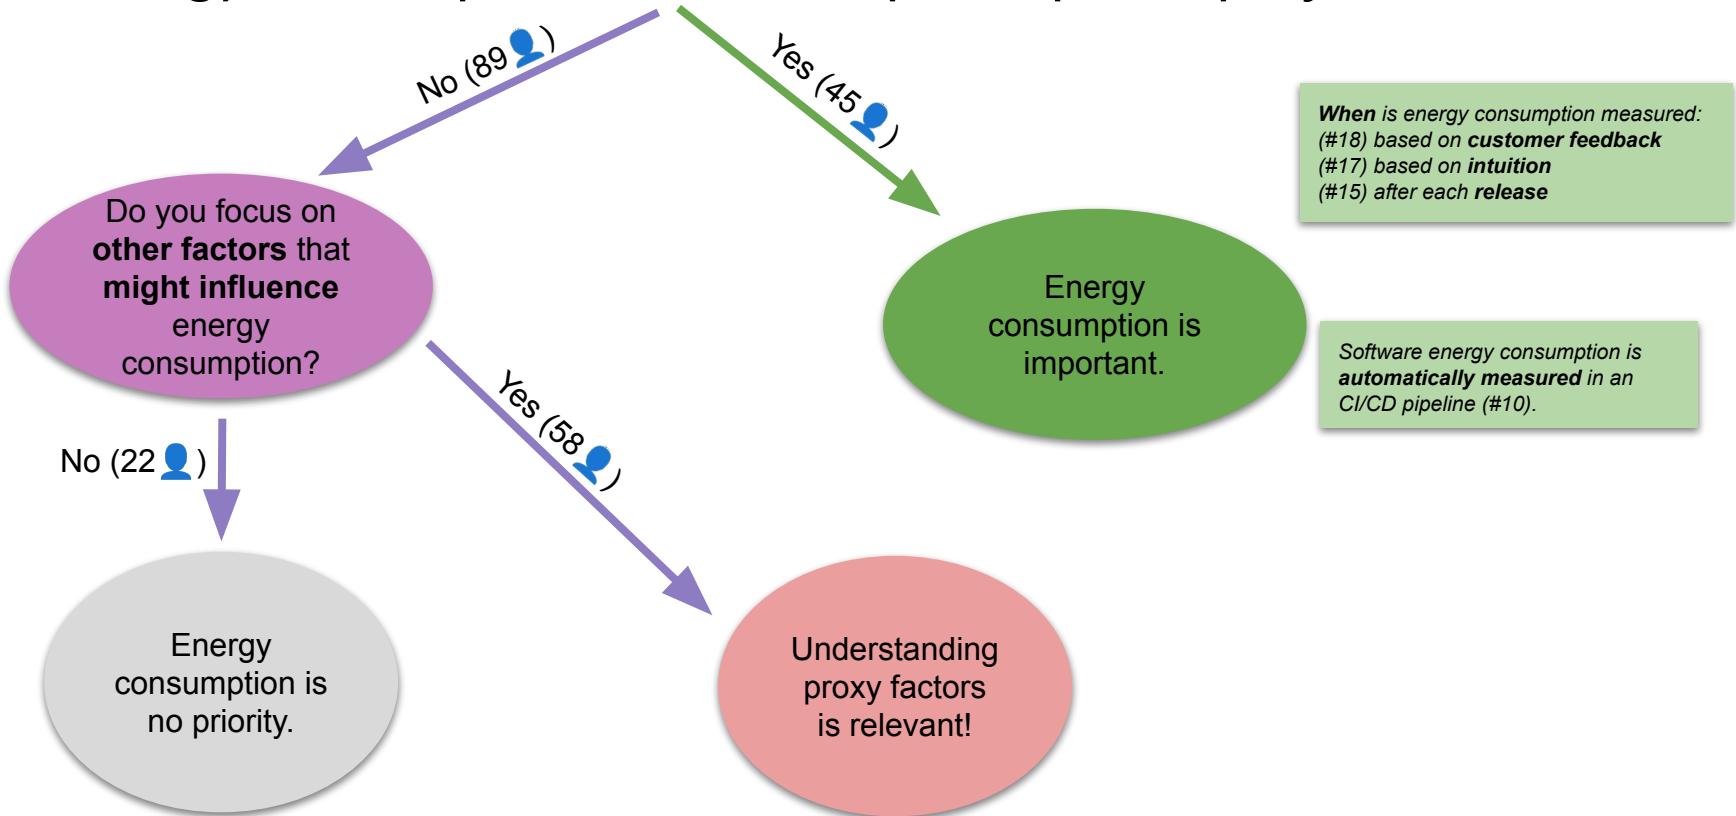
Study on Energy-Conscious Software Development

Energy efficiency in software development is increasingly critical for achieving energy-saving goals and reducing CO₂ emissions. However, it remains unclear to what extent energy consumption is actively considered in software development practices. This survey seeks to capture your experiences with green software development — software design, development, and implementation with a focus on energy efficiency.

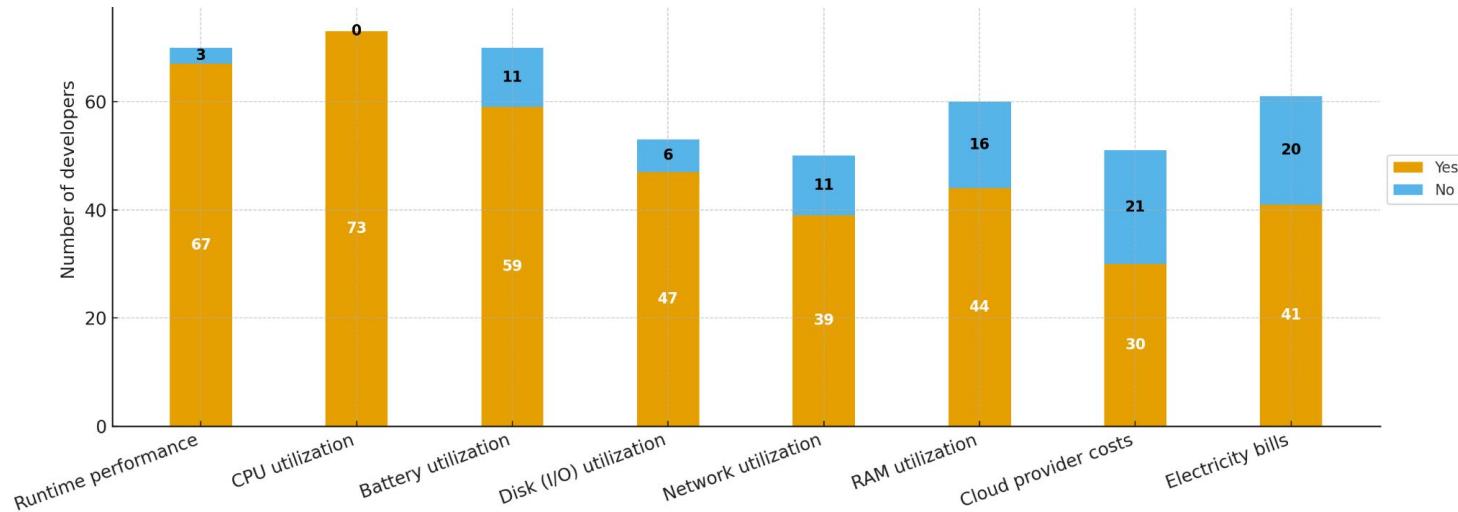
The primary goal of this study is to uncover the current priorities, challenges, and requirements for creating energy-efficient software. Your input is essential for identifying gaps in education, technical adoption, research, tooling, and awareness. By contributing, you help pave the way for a more energy-conscious approach to software engineering in the future.



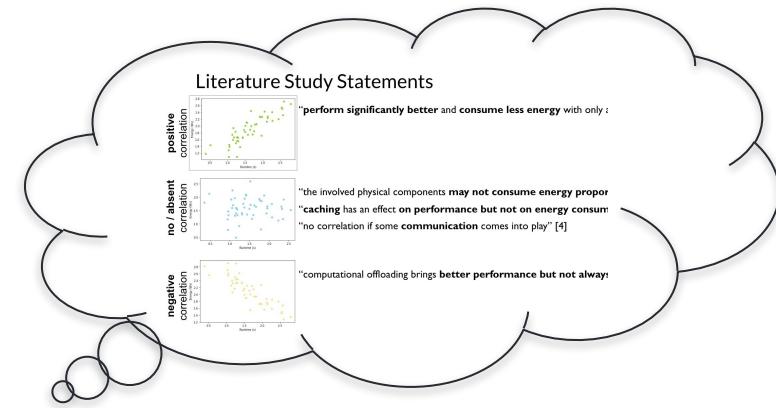
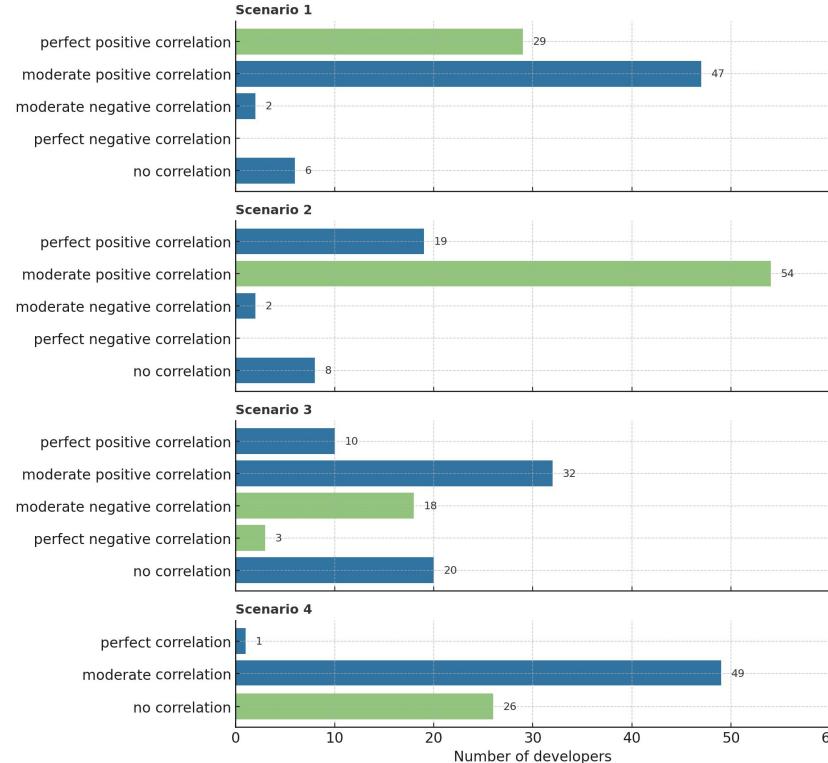
Is energy consumption relevant in participants' projects?



What metrics are considered good proxies for energy?



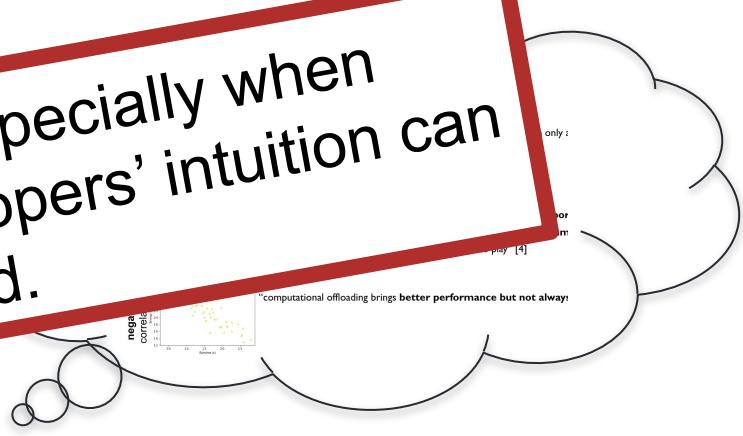
Do Developers Intuitively Know when Performance is a good Proxy for Energy?



Do Developers Intuitively Know when Performance is a good Proxy for Energy?



Finding: It depends! Especially when positively correlated developers' intuition can be trusted.



How to Support Practitioners

Barriers to Overcome

Barriers preventing energy measurement:

- (55) Management is not interested
- (46) Customers are not interested
- (36) Not applicable in our use case

Persistent obstacles:

- (14) Lack of energy measurement tooling
- (10) Energy efficiency is perceived as complex
- (06) Lack of education and training

Contra there are problems:

- (11) Energy consumption is not important
- (06) No perceived challenges
- (05) Energy is not considered a cost driver

Actions

(13) Increasing awareness

(08) Educate developers and stakeholders

(17) Provide better energy measurement tools

(06) 'just' use efficient languages (C, C++, Rust)

(06) Introduce public regulations

(04) Enforce transparency, especially for cloud providers

(05) Adopt efficient algorithms and implementations



UNIVERSITÄT
LEIPZIG

on the job market
from march '26



Feel free to contact me!

max.weber@cs.uni-leipzig.de

<https://github.com/AI-4-SE>

linkedin.com/in/max-weber-84a83415b